

# ASDS-6301 Final Project

## A Regression-Based Analysis of NBA Team Performance

Phuong Trinh, Sashi Kumar Soni, Ahmet Toure

2025-12-09

### Import Packages

```
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr     1.1.4     v readr     2.1.5
## vforcats   1.0.0     v stringr   1.5.1
## v ggplot2   4.0.0     v tibble    3.2.1
## v lubridate 1.9.4     v tidyr    1.3.1
## v purrr    1.0.4
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()   masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

library(lubridate)
library(GGally)
library(glmnet)

## Loading required package: Matrix
##
## Attaching package: 'Matrix'
##
## The following objects are masked from 'package:tidyverse':
## 
##     expand, pack, unpack
##
## Loaded glmnet 4.1-10

library(lme4)
library(broom)
library(DescTools)
library(janitor)

## 
## Attaching package: 'janitor'
```

```

##  

## The following objects are masked from 'package:stats':  

##  

##     chisq.test, fisher.test  

library(dplyr)  

library(tidyr)  

library(ggplot2)  

library(purrr)

```

## Load dataset

```

set.seed(123)  

nba <- read_csv("NBA_FULL_DATASET_2023_2025-2.csv", show_col_types = FALSE)  

head(nba)

```

```

## # A tibble: 6 x 42
##   Team_ID Game_ID GAME_DATE MATCHUP WL      W L W_PCT MIN FGM FGA
##   <dbl> <chr>    <date>   <chr>   <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 1.61e9 002230~ 2023-10-25 ATL @ ~ L      0 1 0 240 39 93
## 2 1.61e9 002230~ 2023-10-27 ATL vs~ L      0 2 0 240 42 87
## 3 1.61e9 002230~ 2023-10-29 ATL @ ~ W      1 2 0.333 240 47 93
## 4 1.61e9 002230~ 2023-10-30 ATL vs~ W      2 2 0.5 240 48 86
## 5 1.61e9 002230~ 2023-11-01 ATL vs~ W      3 2 0.6 240 46 92
## 6 1.61e9 002230~ 2023-11-04 ATL @ ~ W      4 2 0.667 240 45 93
## # i 31 more variables: FG_PCT <dbl>, FG3M <dbl>, FG3A <dbl>, FG3_PCT <dbl>,
## #   FTM <dbl>, FTA <dbl>, FT_PCT <dbl>, OREB <dbl>, DREB <dbl>, REB <dbl>,
## #   AST <dbl>, STL <dbl>, BLK <dbl>, TOV <dbl>, PF <dbl>, PTS <dbl>,
## #   TEAM_NAME <chr>, TEAM_ID <dbl>, SEASON <chr>, OPP_ABBR <chr>,
## #   OPPONENT <chr>, OPPONENT_ID <dbl>, HOME_AWAY <chr>, PREV_GAME_DATE <date>,
## #   DAYS_REST <dbl>, BACK_TO_BACK <dbl>, ARENA_LAT <dbl>, ARENA_LON <dbl>,
## #   PREV_LAT <dbl>, PREV_LON <dbl>, TRAVEL_DISTANCE <dbl>

```

## Data Preprocessing

### Data Cleaning

```

# Check variable names
nba <- nba %>% clean_names()

```

### Compute PLUS/MINUS

```

library(dplyr)
pm <- nba %>%

```

```

select(game_id, team_name, opponent, pts) %>%
inner_join(
  nba %>% select(game_id, team_name, pts) %>%
    rename(opp_team = team_name, opp_pts = pts),
    by = "game_id"
) %>%
filter(team_name != opp_team) %>%
group_by(game_id, team_name, opponent) %>%
summarise(plus_minus = first(pts - opp_pts), .groups = "drop")

## Warning in inner_join(., nba %>% select(game_id, team_name, pts) %>% rename(opp_team = team_name, :
## i Row 1 of 'x' matches multiple rows in 'y'.
## i Row 16 of 'y' matches multiple rows in 'x'.
## i If a many-to-many relationship is expected, set 'relationship =
##   "many-to-many"' to silence this warning.

nba <- nba %>%
  left_join(pm, by = c("game_id", "team_name", "opponent"))

summary(nba$plus_minus)

##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
##     -62     -10      0       0      10      62

# Check missing values
missing_summary <- colSums(is.na(nba))
missing_summary

##          team_id      game_id      game_date      matchup        wl
##             0           0           0           0           0
##            w           1         w_pct        min        fgm
##            0           0           0           0           0
##           fga        fg_pct      fg3m        fg3a fg3_pct
##            0           0           0           0           0
##           ftm        fta_pct      ft_pct      oreb      dreb
##            0           0           0           0           0
##           reb        ast_pct      stl_pct      blk      tov
##            0           0           0           0           0
##           pf         pts_team team_name team_id_2      season
##            0           0           0           0           0
##      opp_abbr      opponent opponent_id home_away prev_game_date
##            0           0           0           0           30
## days_rest back_to_back arena_lat arena_lon      prev_lat
##            0           0           0           0           30
## prev_lon travel_distance plus_minus
##            30           0           0
missing_summary[missing_summary > 0]

## prev_game_date      prev_lat      prev_lon
##            30           30           30

```

```

## Handle Missing Values

# 1. Replace missing travel distance with 0 (first game of each team)
nba <- nba %>%
  mutate(
    travel_distance = ifelse(is.na(travel_distance), 0, travel_distance)
  )

# 2. Replace missing previous coordinates with current coordinates
nba <- nba %>%
  mutate(
    prev_lat = ifelse(is.na(prev_lat), arena_lat, prev_lat),
    prev_lon = ifelse(is.na(prev_lon), arena_lon, prev_lon)
  )

# 3. Replace missing previous game date with the season start date
# (or leave as NA - either is acceptable)
nba <- nba %>%
  mutate(
    prev_game_date = ifelse(is.na(prev_game_date), game_date, prev_game_date)
  )

# Check again
colSums(is.na(nba))

```

##	team_id	game_id	game_date	matchup	wl
##	0	0	0	0	0
##	w	1	w_pct	min	fgm
##	0	0	0	0	0
##	fga	fg_pct	fg3m	fg3a	fg3_pct
##	0	0	0	0	0
##	ftm	fta	ft_pct	oreb	dreb
##	0	0	0	0	0
##	reb	ast	stl	blk	tov
##	0	0	0	0	0
##	pf	pts	team_name	team_id_2	season
##	0	0	0	0	0
##	opp_abbr	opponent	opponent_id	home_away	prev_game_date
##	0	0	0	0	0
##	days_rest	back_to_back	arena_lat	arena_lon	prev_lat
##	0	0	0	0	0
##	prev_lon	travel_distance	plus_minus		
##	0	0	0		

```

# Check duplicate records
num_duplicates <- sum(duplicated(nba))
num_duplicates

```

```
## [1] 0
```

```

# Remove duplicates
nba <- nba %>% distinct()

```

```
# Drop team_id and game_id  
nba <- nba %>% select(-team_id, -game_id)
```

## Descriptive Statistic

```
# Check data types  
str(nba)
```

```
## # tibble [4,920 x 41] (S3: tbl_df/tbl/data.frame)
## # $ game_date      : Date[1:4920], format: "2023-10-25" "2023-10-27" ...
## # $ matchup        : chr [1:4920] "ATL @ CHA" "ATL vs. NYK" "ATL @ MIL" "ATL vs. MIN" ...
## # $ wl             : chr [1:4920] "L" "L" "W" "W" ...
## # $ w              : num [1:4920] 0 0 1 2 3 4 4 5 5 6 ...
## # $ l              : num [1:4920] 1 2 2 2 2 2 3 3 4 4 ...
## # $ w_pct          : num [1:4920] 0 0 0.333 0.5 0.6 0.667 0.571 0.625 0.556 0.6 ...
## # $ min            : num [1:4920] 240 240 240 240 240 240 240 240 240 240 ...
## # $ fgm            : num [1:4920] 39 42 47 48 46 45 38 41 38 45 ...
## # $ fga            : num [1:4920] 93 87 93 86 92 93 102 85 87 87 ...
## # $ fg_pct         : num [1:4920] 0.419 0.483 0.505 0.558 0.5 0.484 0.373 0.482 0.437 0.517 ...
## # $ fg3m           : num [1:4920] 5 12 15 14 9 14 14 15 18 12 ...
## # $ fg3a           : num [1:4920] 29 32 37 30 32 41 42 39 46 33 ...
## # $ fg3_pct        : num [1:4920] 0.172 0.375 0.405 0.467 0.281 0.341 0.333 0.385 0.391 0.364 ...
## # $ ftm            : num [1:4920] 27 24 18 17 29 19 27 23 15 24 ...
## # $ fta            : num [1:4920] 33 30 22 18 32 24 32 30 18 26 ...
## # $ ft_pct         : num [1:4920] 0.818 0.8 0.818 0.944 0.906 0.792 0.844 0.767 0.833 0.923 ...
## # $ oreb           : num [1:4920] 12 9 13 4 14 19 25 8 15 10 ...
## # $ dreb           : num [1:4920] 30 35 33 32 43 33 34 28 33 31 ...
## # $ reb             : num [1:4920] 42 44 46 36 57 52 59 36 48 41 ...
## # $ ast             : num [1:4920] 24 28 32 28 26 28 27 22 23 28 ...
## # $ stl             : num [1:4920] 12 7 15 6 8 4 8 13 6 9 ...
## # $ blk             : num [1:4920] 1 6 2 7 3 6 3 7 6 7 ...
## # $ tov             : num [1:4920] 12 14 17 9 20 11 15 17 21 14 ...
## # $ pf               : num [1:4920] 19 20 17 12 16 16 16 25 19 19 ...
## # $ pts             : num [1:4920] 110 120 127 127 130 123 117 120 109 126 ...
## # $ team_name       : chr [1:4920] "Atlanta Hawks" "Atlanta Hawks" "Atlanta Hawks" "Atlanta Hawks" ...
## # $ team_id_2        : num [1:4920] 1.61e+09 1.61e+09 1.61e+09 1.61e+09 1.61e+09 ...
## # $ season          : chr [1:4920] "2023-24" "2023-24" "2023-24" "2023-24" ...
## # $ opp_abbr        : chr [1:4920] "CHA" "NYK" "MIL" "MIN" ...
## # $ opponent         : chr [1:4920] "Charlotte Hornets" "New York Knicks" "Milwaukee Bucks" "Minnesota ...
## # $ opponent_id      : num [1:4920] 1.61e+09 1.61e+09 1.61e+09 1.61e+09 1.61e+09 ...
## # $ home_away        : chr [1:4920] "Away" "Home" "Away" "Home" ...
## # $ prev_game_date  : num [1:4920] 19655 19655 19657 19659 19660 ...
## # $ days_rest        : num [1:4920] 5 2 2 1 2 3 2 3 2 3 ...
## # $ back_to_back     : num [1:4920] 0 0 0 1 0 0 0 0 0 0 ...
## # $ arena_lat        : num [1:4920] 35.2 33.8 43 33.8 33.8 ...
## # $ arena_lon        : num [1:4920] -80.8 -84.4 -87.9 -84.4 -84.4 ...
## # $ prev_lat          : num [1:4920] 35.2 35.2 33.8 43 33.8 ...
## # $ prev_lon          : num [1:4920] -80.8 -80.8 -84.4 -87.9 -84.4 ...
## # $ travel_distance: num [1:4920] 0 227 668 668 0 ...
## # $ plus_minus       : num [1:4920] -6 -6 17 14 9 18 -9 1 -8 6 ...
```

```

# Convert categorical variables to factors
nba <- nba %>%
  mutate(
    game_date      = as_date(game_date),
    prev_game_date = as_date(prev_game_date),
    season         = factor(season),
    wl             = factor(wl, levels = c("L", "W")),
    home_away      = factor(home_away),
    back_to_back   = factor(back_to_back, levels = c("0", "1")),
    team_name      = factor(team_name),
    opponent        = factor(opponent)
  )

summary(nba)

```

```

##      game_date           matchup          wl          w
##  Min.   :2023-10-24  Length:4920    L:2460  Min.   : 0.00
##  1st Qu.:2024-01-19  Class  :character  W:2460  1st Qu.: 9.00
##  Median :2024-07-18  Mode   :character            Median :19.00
##  Mean   :2024-07-19                    Mean   :20.74
##  3rd Qu.:2025-01-18                    3rd Qu.:31.00
##  Max.   :2025-04-13                    Max.   :68.00
##
##      l           w_pct          min          fgm
##  Min.   : 0.00  Min.   :0.0000  Min.   :240.0  Min.   :22.00
##  1st Qu.: 9.00  1st Qu.:0.3850  1st Qu.:240.0  1st Qu.:38.00
##  Median :19.00  Median :0.5220  Median :240.0  Median :42.00
##  Mean   :20.76  Mean   :0.4998  Mean   :241.3  Mean   :41.93
##  3rd Qu.:30.00  3rd Qu.:0.6110  3rd Qu.:240.0  3rd Qu.:45.00
##  Max.   :68.00  Max.   :1.0000  Max.   :290.0  Max.   :65.00
##
##      fga          fg_pct          fg3m          fg3a
##  Min.   : 67.00  Min.   :0.2770  Min.   : 2.00  Min.   :12.00
##  1st Qu.: 84.00  1st Qu.:0.4340  1st Qu.:10.00  1st Qu.:31.00
##  Median : 89.00  Median :0.4710  Median :13.00  Median :36.00
##  Mean   : 89.07  Mean   :0.4717  Mean   :13.19  Mean   :36.34
##  3rd Qu.: 94.00  3rd Qu.:0.5100  3rd Qu.:16.00  3rd Qu.:41.00
##  Max.   :119.00  Max.   :0.6890  Max.   :29.00  Max.   :63.00
##
##      fg3_pct          ftm          fta          ft_pct
##  Min.   :0.0690  Min.   : 0.00  Min.   : 0.0  Min.   :0.0000
##  1st Qu.:0.3080  1st Qu.:13.00  1st Qu.:17.0  1st Qu.:0.7140
##  Median :0.3610  Median :17.00  Median :21.0  Median :0.7890
##  Mean   :0.3624  Mean   :16.97  Mean   :21.7  Mean   :0.7814
##  3rd Qu.:0.4150  3rd Qu.:21.00  3rd Qu.:26.0  3rd Qu.:0.8500
##  Max.   :0.6800  Max.   :44.00  Max.   :53.0  Max.   :1.0000
##
##      oreb          dreb          reb          ast
##  Min.   : 0.00  Min.   :16.00  Min.   :23.00  Min.   : 9.00
##  1st Qu.: 8.00  1st Qu.:29.00  1st Qu.:39.00  1st Qu.:23.00
##  Median :11.00  Median :33.00  Median :44.00  Median :26.00
##  Mean   :10.84  Mean   :32.99  Mean   :43.83  Mean   :26.61
##  3rd Qu.:13.00  3rd Qu.:37.00  3rd Qu.:48.00  3rd Qu.:30.00

```

```

##   Max.    :28.00   Max.    :55.00   Max.    :74.00   Max.    :50.00
##
##          stl            blk            tov            pf            pts
##  Min.    : 0.00   Min.    :0.0000   Min.    : 2.0   Min.    : 4.00   Min.    : 67
##  1st Qu.: 6.00   1st Qu.: 3.0000   1st Qu.:10.0   1st Qu.:16.00   1st Qu.:105
##  Median : 8.00   Median : 5.0000   Median :13.0   Median :18.00   Median :114
##  Mean   : 7.84   Mean   : 5.0090   Mean   :13.2   Mean   :18.66   Mean   :114
##  3rd Qu.:10.00   3rd Qu.: 7.0000   3rd Qu.:16.0   3rd Qu.:21.00   3rd Qu.:123
##  Max.   :22.00   Max.   :18.0000   Max.   :31.0   Max.   :35.00   Max.   :162
##
##          team_name      team_id_2      season
##  Atlanta Hawks     : 164   Min.    :1.611e+09  2023-24:2460
##  Boston Celtics    : 164   1st Qu.:1.611e+09  2024-25:2460
##  Brooklyn Nets     : 164   Median :1.611e+09
##  Charlotte Hornets : 164   Mean    :1.611e+09
##  Chicago Bulls     : 164   3rd Qu.:1.611e+09
##  Cleveland Cavaliers: 164   Max.   :1.611e+09
##  (Other)           :3936
##
##          opp_abbr        opponent      opponent_id      home_away
##  Length:4920       Atlanta Hawks   : 164   Min.    :1.611e+09  Away:2460
##  Class :character  Boston Celtics   : 164   1st Qu.:1.611e+09  Home:2460
##  Mode  :character  Brooklyn Nets   : 164   Median :1.611e+09
##                      Charlotte Hornets : 164   Mean   :1.611e+09
##                      Chicago Bulls    : 164   3rd Qu.:1.611e+09
##                      Cleveland Cavaliers: 164   Max.   :1.611e+09
##  (Other)           :3936
##
##          prev_game_date    days_rest    back_to_back    arena_lat
##  Min.    :2023-10-24   Min.    : 1.000   0:4048      Min.    :25.78
##  1st Qu.:2024-01-17   1st Qu.: 2.000   1: 872      1st Qu.:33.76
##  Median :2024-04-13   Median : 2.000                Median :38.90
##  Mean   :2024-07-16   Mean   : 3.299                Mean   :37.32
##  3rd Qu.:2025-01-16   3rd Qu.: 2.000                3rd Qu.:41.50
##  Max.   :2025-04-11   Max.   :193.000               Max.   :45.53
##
##          arena_lon      prev_lat      prev_lon      travel_distance
##  Min.    :-122.67   Min.    :25.78   Min.    :-122.67   Min.    : 0.0
##  1st Qu.:-105.01   1st Qu.:33.76   1st Qu.:-105.01   1st Qu.: 0.0
##  Median :-88.98    Median :38.90    Median :-87.92    Median : 420.5
##  Mean   :-93.27    Mean   :37.32    Mean   :-93.26    Mean   : 525.1
##  3rd Qu.:-80.84   3rd Qu.:41.50   3rd Qu.:-80.84   3rd Qu.: 857.9
##  Max.   :-71.06   Max.   :45.53    Max.   :-71.06   Max.   :2707.4
##
##          plus_minus
##  Min.   :-62
##  1st Qu.:-10
##  Median : 0
##  Mean   : 0
##  3rd Qu.: 10
##  Max.   : 62
##

```

## Feature Engineering

```
# Rest_days already computed
summary(nba$days_rest)

##      Min. 1st Qu. Median     Mean 3rd Qu.    Max.
##      1.000   2.000   2.000   3.299   2.000 193.000

# Back_to_back already computed
table(nba$back_to_back)

##
##      0      1
## 4048  872

#Travel_distance already computed
summary(nba$travel_distance)

##      Min. 1st Qu. Median     Mean 3rd Qu.    Max.
##      0.0      0.0   420.5   525.1   857.9 2707.4

# Home_Away convert to numeric
nba <- nba %>%
  mutate(home_indicator = if_else(home_away == "Home", 1, 0))

# Opponent_strength (opp win percentage to date)
nba <- nba %>%
  group_by(opponent) %>%
  arrange(game_date) %>%
  mutate(
    opponent_strength = lag(cumsum(w == 1) / row_number(), default = 0)
  ) %>%
  ungroup()

summary(nba$opponent_strength)

##      Min. 1st Qu. Median     Mean 3rd Qu.    Max.
## 0.00000 0.02500 0.03750 0.06403 0.05556 1.00000
```

## Outlier Detection

```
# IQR method
# travel_distance
Q1 <- quantile(nba$travel_distance, 0.25, na.rm = TRUE)
Q3 <- quantile(nba$travel_distance, 0.75, na.rm = TRUE)
IQR_val <- Q3 - Q1

upper_bound <- Q3 + 1.5 * IQR_val
```

```

lower_bound <- Q1 - 1.5 * IQR_val

# days_rest
Q1_rest <- quantile(nba$days_rest, 0.25, na.rm = TRUE)
Q3_rest <- quantile(nba$days_rest, 0.75, na.rm = TRUE)
IQR_rest <- Q3_rest - Q1_rest

upper_rest <- Q3_rest + 1.5 * IQR_rest
lower_rest <- Q1_rest - 1.5 * IQR_rest

# plus_minus
Q1_pm <- quantile(nba$plus_minus, 0.25, na.rm = TRUE)
Q3_pm <- quantile(nba$plus_minus, 0.75, na.rm = TRUE)
IQR_pm <- Q3_pm - Q1_pm

upper_pm <- Q3_pm + 1.5 * IQR_pm
lower_pm <- Q1_pm - 1.5 * IQR_pm

# Opponent strength
Q1_str <- quantile(nba$opponent_strength, 0.25, na.rm = TRUE)
Q3_str <- quantile(nba$opponent_strength, 0.75, na.rm = TRUE)
IQR_str <- Q3_str - Q1_str

upper_str <- Q3_str + 1.5 * IQR_str
lower_str <- Q1_str - 1.5 * IQR_str

# Table summary
outlier_summary <- tibble(
  variable = c("travel_distance", "days_rest", "plus_minus", "opponent_strength"),
  lower_bound = c(lower_bound, lower_rest, lower_pm, lower_str),
  upper_bound = c(upper_bound, upper_rest, upper_pm, upper_str)
)

outlier_summary

```

```

## # A tibble: 4 x 3
##   variable      lower_bound  upper_bound
##   <chr>          <dbl>        <dbl>
## 1 travel_distance -1287.       2145.
## 2 days_rest         2           2
## 3 plus_minus        -40          40
## 4 opponent_strength -0.0208     0.101

```

```

# Visualize the outliers
library(tidyr)
library(ggplot2)

plot_data <- nba %>%
  select(travel_distance, days_rest, plus_minus, opponent_strength) %>%
  pivot_longer(cols = everything(), names_to = "variable", values_to = "value")

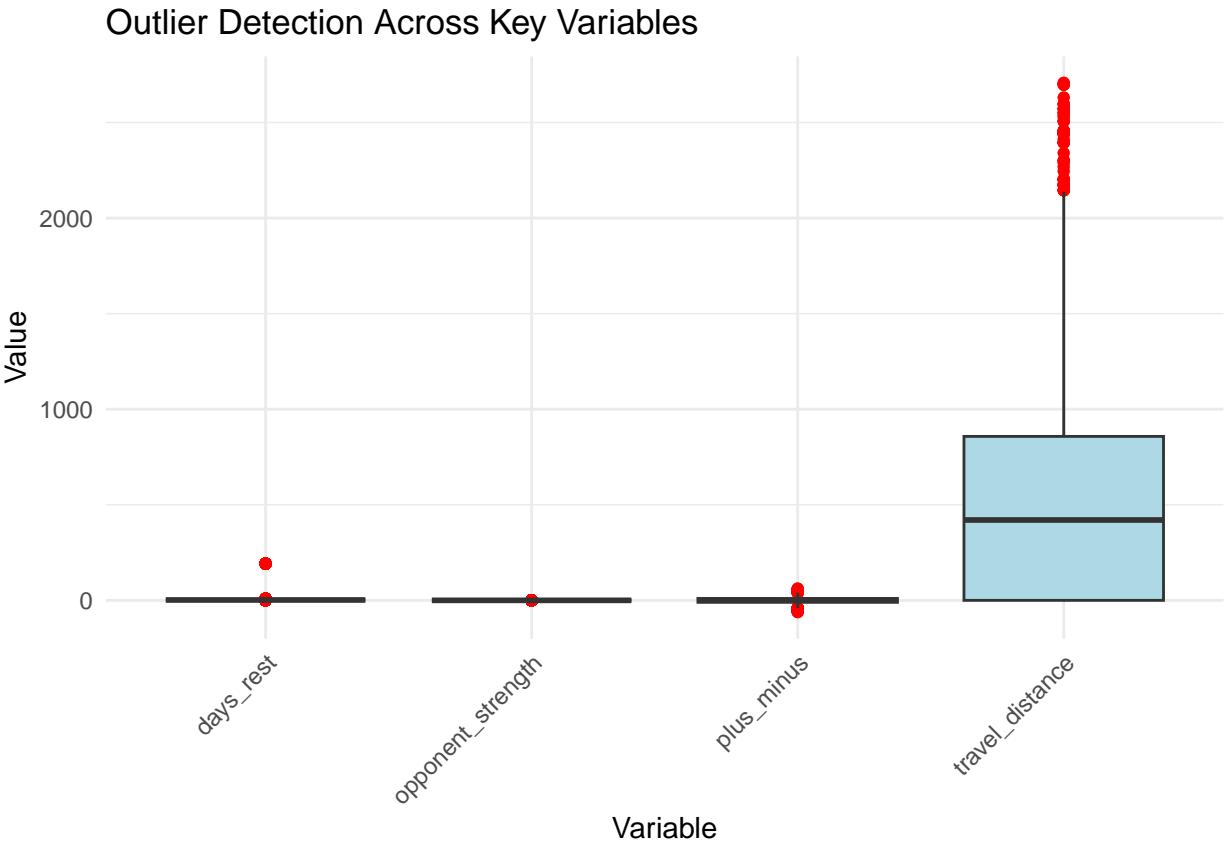
ggplot(plot_data, aes(x = variable, y = value)) +
  geom_boxplot(outlier.color = "red", fill = "lightblue") +

```

```

theme_minimal() +
labs(title = "Outlier Detection Across Key Variables",
x = "Variable", y = "Value") +
theme(axis.text.x = element_text(angle = 45, hjust = 1))

```



```

# Handling Outliers by Winsorization
winsorize_manual <- function(x, p = c(0.01, 0.99)) {
  qnt <- quantile(x, probs = p, na.rm = TRUE)
  x[x < qnt[1]] <- qnt[1]
  x[x > qnt[2]] <- qnt[2]
  return(x)
}

nba <- nba %>%
  mutate(
    travel_distance_w = winsorize_manual(travel_distance)
  )

library(tidyr)

plot_df <- nba %>%
  select(travel_distance, travel_distance_w) %>%
  pivot_longer(cols = everything(), names_to = "type", values_to = "value")

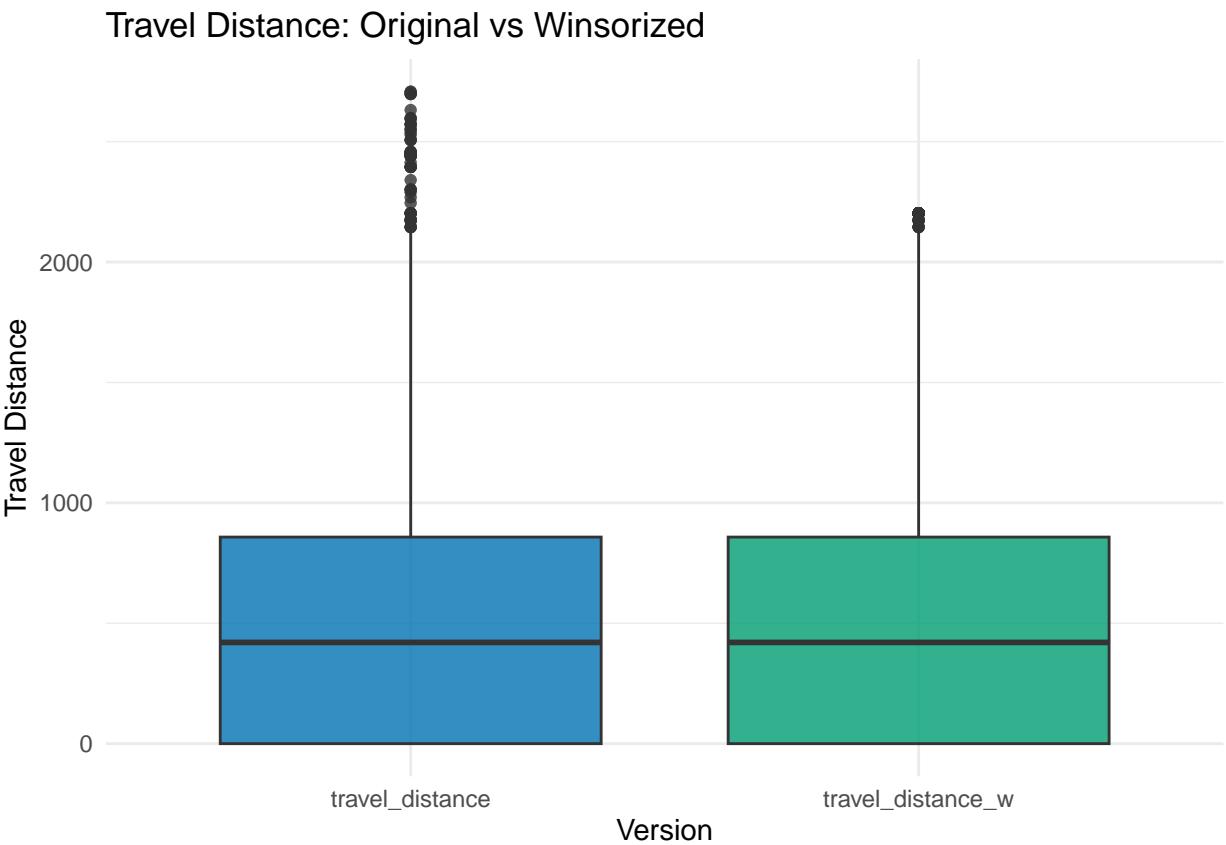
ggplot(plot_df, aes(x = type, y = value)) +

```

```

geom_boxplot(fill = c("#0072B2", "#009E73"), alpha = 0.8) +
  labs(
    title = "Travel Distance: Original vs Winsorized",
    x = "Version",
    y = "Travel Distance"
  ) +
  theme_minimal()

```



## Normalization and Scaling

```

nba <- nba %>%
  mutate(
    scale_travel  = scale(travel_distance),
    scale_rest    = scale(days_rest),
    scale_plus    = scale(plus_minus),
    scale_opp_str = scale(opponent_strength)
  )

## View final cleaned dataset
#glimpse(nba)

```

```

nba <- nba %>%
  mutate(
    back_to_back = as.character(back_to_back),
    back_to_back_num = ifelse(back_to_back == "1", 1L, 0L),
    travel_distance_w = ifelse(is.na(travel_distance_w), travel_distance, travel_distance_w)
  )

# Check class distribution for the target variable 'w'

table(nba$wl)

##          L          W
## 2460 2460

prop.table(table(nba$wl))

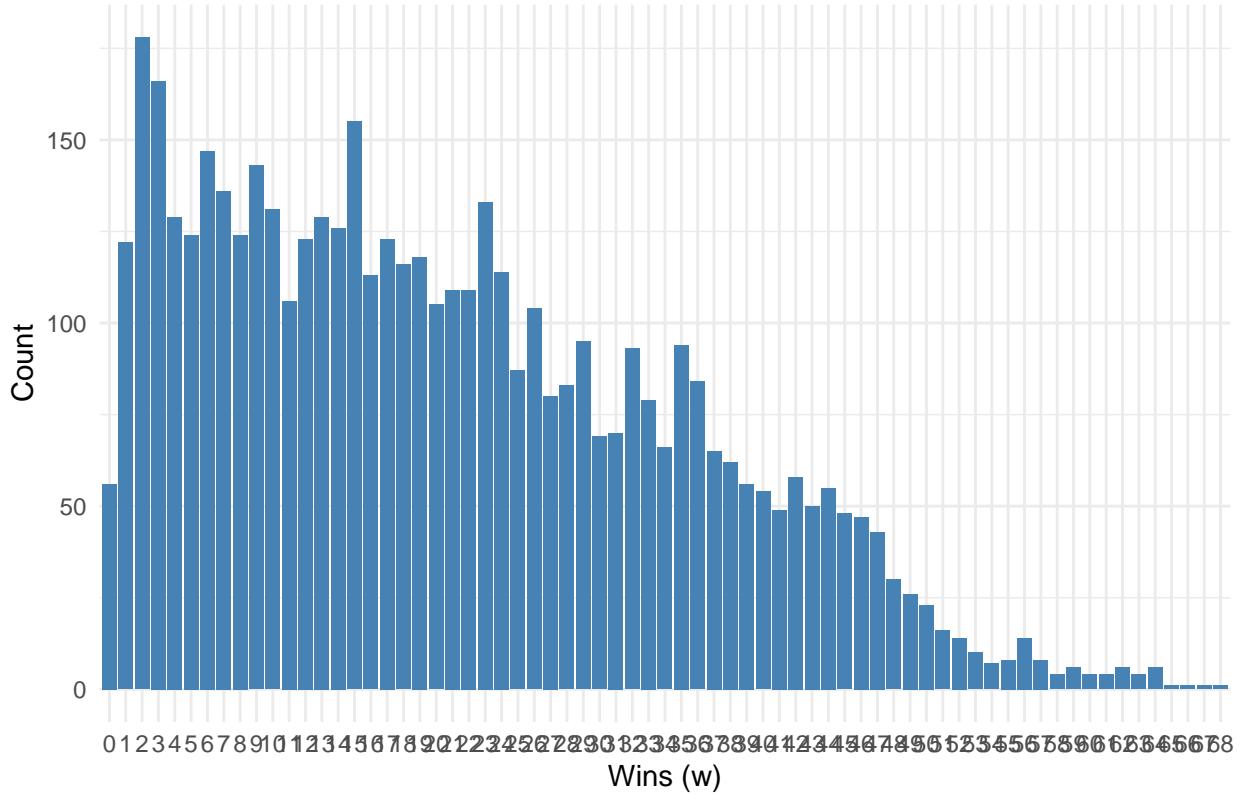
##          L          W
## 0.5 0.5

# Visual bar plot
library(ggplot2)

ggplot(nba, aes(x = factor(w))) +
  geom_bar(fill = "steelblue") +
  labs(title = "Distribution of Target Variable: Wins (w)",
       x = "Wins (w)",
       y = "Count") +
  theme_minimal()

```

## Distribution of Target Variable: Wins (w)



## EDA

### Trend over Time: B2B Games per Season

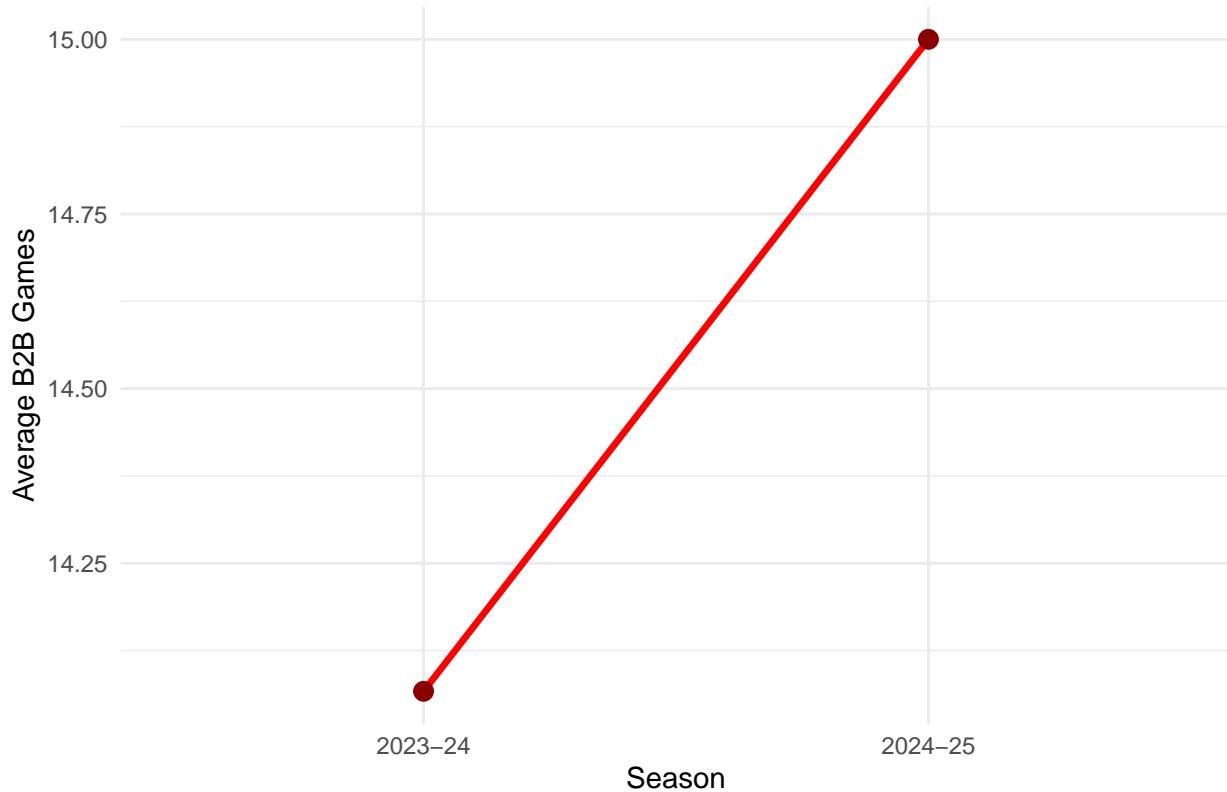
NBA scheduling has reduced B2B games over time to improve player rest.

```
library(dplyr)
library(ggplot2)

b2b_trend <- nba %>%
  group_by(season, team_name) %>%
  summarise(total_b2b = sum(back_to_back_num, na.rm = TRUE), .groups = "drop") %>%
  group_by(season) %>%
  summarise(avg_b2b = base::mean(total_b2b, na.rm = TRUE), .groups = "drop")

ggplot(b2b_trend, aes(season, avg_b2b, group = 1)) +
  geom_line(color = "red", linewidth = 1.2) +
  geom_point(color = "darkred", size = 3) +
  theme_minimal() +
  labs(title = "Trend in Average Back-to-Back Games per Team",
       x = "Season", y = "Average B2B Games")
```

## Trend in Average Back-to-Back Games per Team

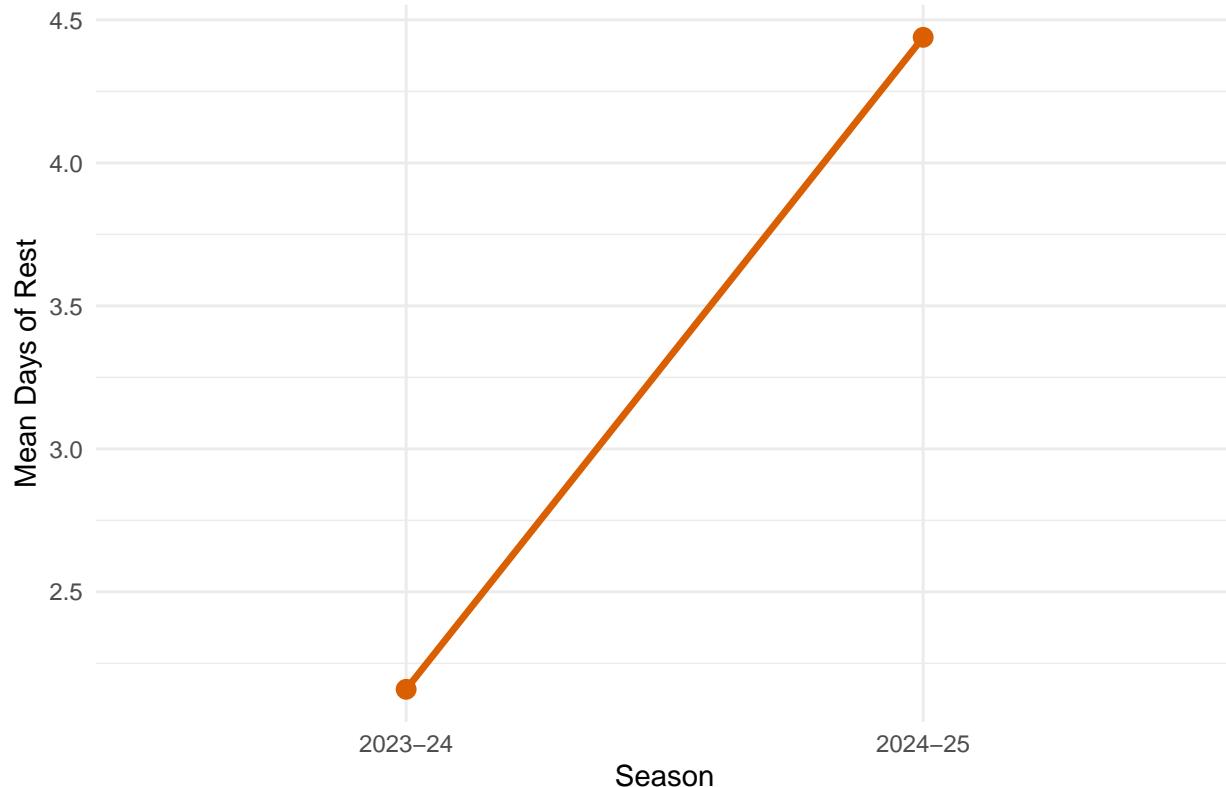


## Average rest days by season

```
rest_trend <- nba %>%
  group_by(season) %>%
  summarise(avg_rest = base::mean(days_rest, na.rm = TRUE), .groups = "drop")

ggplot(rest_trend, aes(season, avg_rest, group = 1)) +
  geom_line(color = "#d95f02", linewidth = 1.2) +
  geom_point(color = "#d95f02", size = 3) +
  theme_minimal() +
  labs(title = "Average Days of Rest by Season",
       x = "Season", y = "Mean Days of Rest")
```

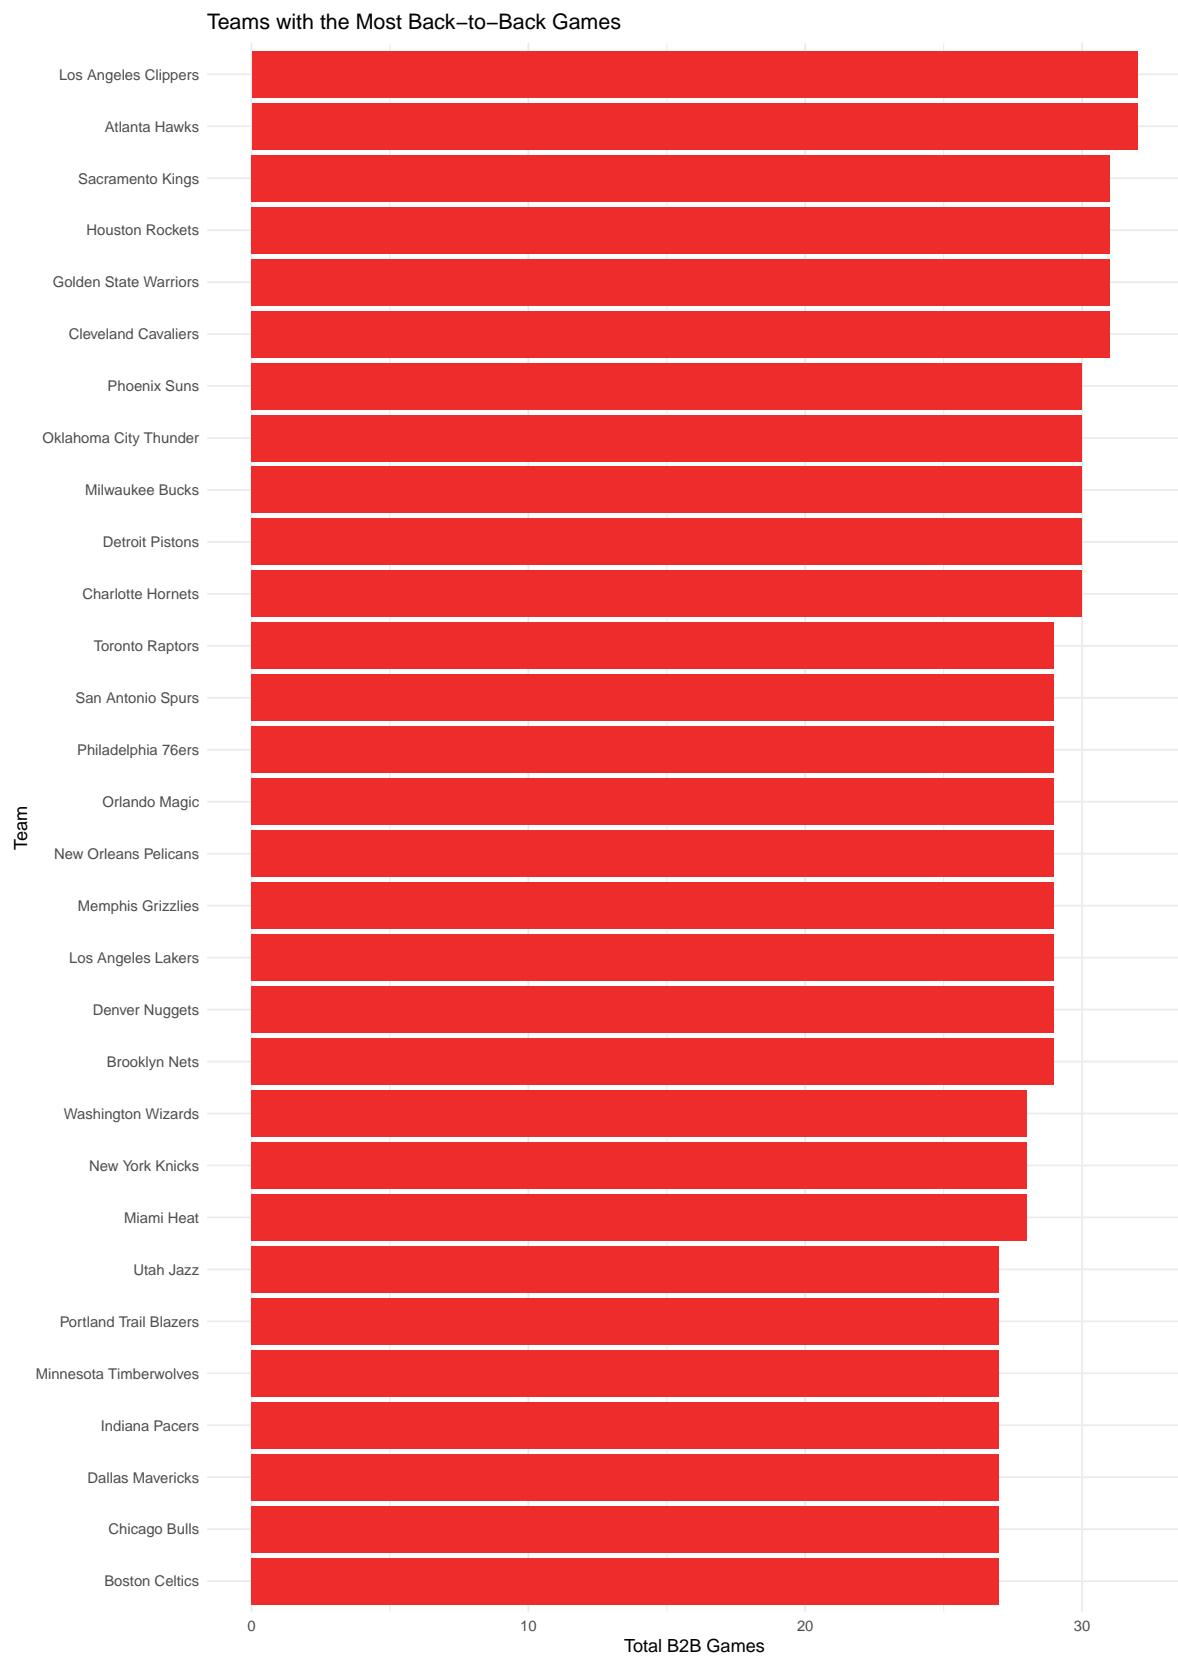
## Average Days of Rest by Season



## By team: Most B2B and most rest

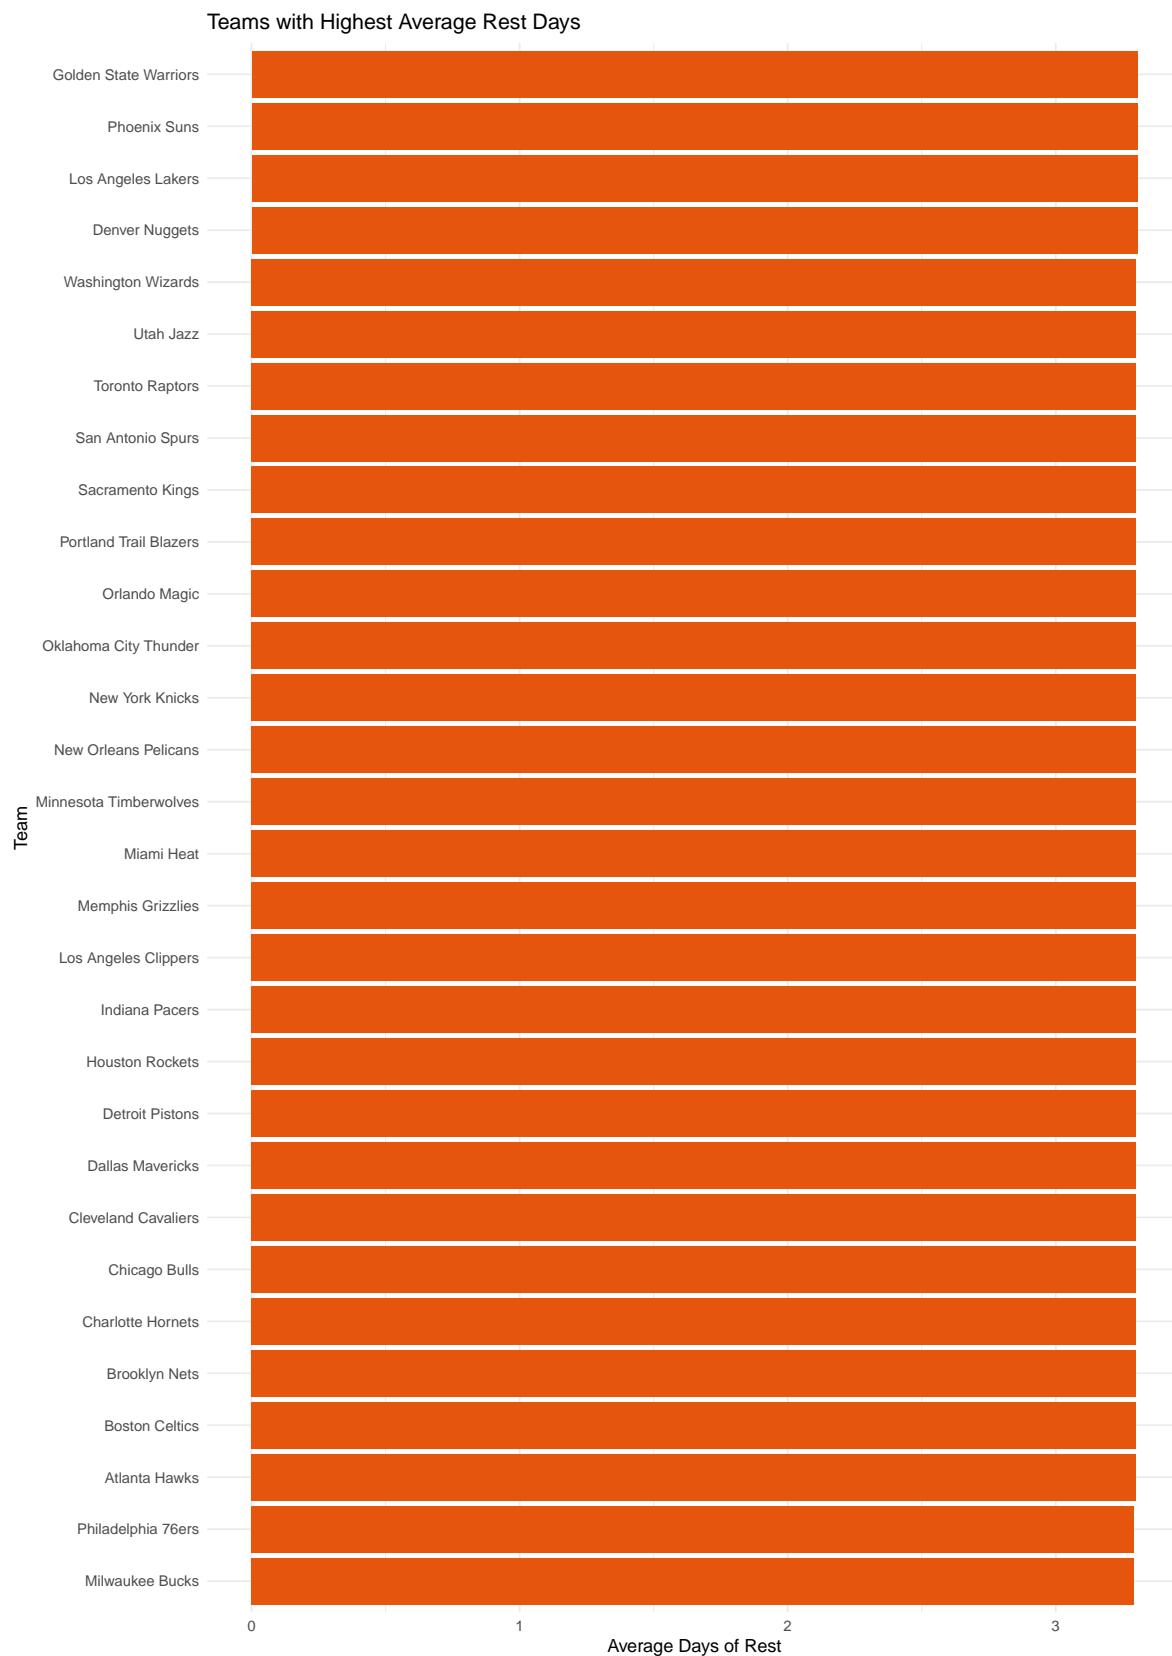
```
# (A) Teams with the most B2B games (all seasons combined)
b2b_team <- nba %>%
  group_by(team_name) %>%
  summarise(total_b2b = sum(back_to_back_num, na.rm = TRUE), .groups = "drop") %>%
  arrange(desc(total_b2b))

ggplot(b2b_team, aes(x = reorder(team_name, total_b2b), y = total_b2b)) +
  geom_col(fill = "firebrick2") +
  coord_flip() +
  theme_minimal() +
  labs(title = "Teams with the Most Back-to-Back Games",
       x = "Team", y = "Total B2B Games")
```



```
# (B) Teams with the highest average rest
rest_team <- nba %>%
  group_by(team_name) %>%
  summarise(avg_rest = base::mean(days_rest, na.rm = TRUE), .groups = "drop") %>%
  arrange(desc(avg_rest))

ggplot(rest_team, aes(x = reorder(team_name, avg_rest), y = avg_rest)) +
  geom_col(fill = "#e6550d") +
  coord_flip() +
  theme_minimal() +
  labs(title = "Teams with Highest Average Rest Days",
       x = "Team", y = "Average Days of Rest")
```



## Interactive full-season schedule viewer

```
library(dplyr)
library(plotly)

## 
## Attaching package: 'plotly'

## The following object is masked from 'package:ggplot2':
## 
##     last_plot

## The following object is masked from 'package:stats':
## 
##     filter

## The following object is masked from 'package:graphics':
## 
##     layout

team_schedule <- function(team) {

  df <- nba %>%
    filter(team_name == team) %>%
    arrange(game_date) %>%
    mutate(
      back_to_back_num = case_when(
        as.character(back_to_back) == "1" ~ 1L,
        TRUE ~ 0L
      ),
      travel_distance_w = ifelse(is.na(travel_distance_w),
                                 travel_distance,
                                 travel_distance_w),
      # jitter y-position so all points aren't on a single line
      y_jitter = runif(n(), -0.3, 0.3),

      # color rules
      color_code = case_when(
        back_to_back_num == 1 ~ "red",
        days_rest >= 3 ~ "darkgreen",
        TRUE ~ "grey"
      ),
      # shape: home/away
      marker_shape = ifelse(home_away == "Home", "circle", "square")
    )

  plot_ly(
    data = df,
    x = ~game_date,
    y = ~y_jitter,
```

```

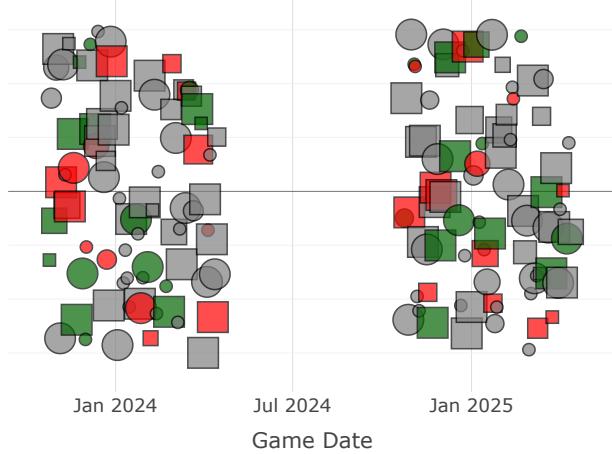
type = "scatter",
mode = "markers",
marker = list(
    size = ~pmin(travel_distance_w / 50 + 8, 20),    # travel scaling
    color = df$color_code,
    symbol = ~marker_shape,
    line = list(width = 1, color = "black")
),
text = ~paste0(
    "<b>Date:</b> ", game_date,
    "<br><b>Opponent:</b> ", opponent,
    "<br><b>Home/Away:</b> ", home_away,
    "<br><b>Days Rest:</b> ", days_rest,
    "<br><b>Travel (mi):</b> ", round(travel_distance_w, 0),
    "<br><b>B2B:</b> ", ifelse(back_to_back_num == 1, "Yes", "No")
),
hoverinfo = "text"
) %>%
layout(
    title = paste0(team, " - Season Schedule Timeline"),
    xaxis = list(
        title = "Game Date",
        showgrid = TRUE,
        gridcolor = "lightgrey"
    ),
    yaxis = list(
        title = "",
        showticklabels = FALSE
    ),
    legend = list(orientation = "h")
)
}

# Example:
team_schedule("Dallas Mavericks")

```

## file:///private/var/folders/2x/m3nyg7x50mzg\_7wsnzcm7hxro000gn/T/Rtmp7CdIv4/file11453366e15f1/widget

## Dallas Mavericks — Season Schedule Timeline



## Correlation Heatmap

```
library(corrplot)

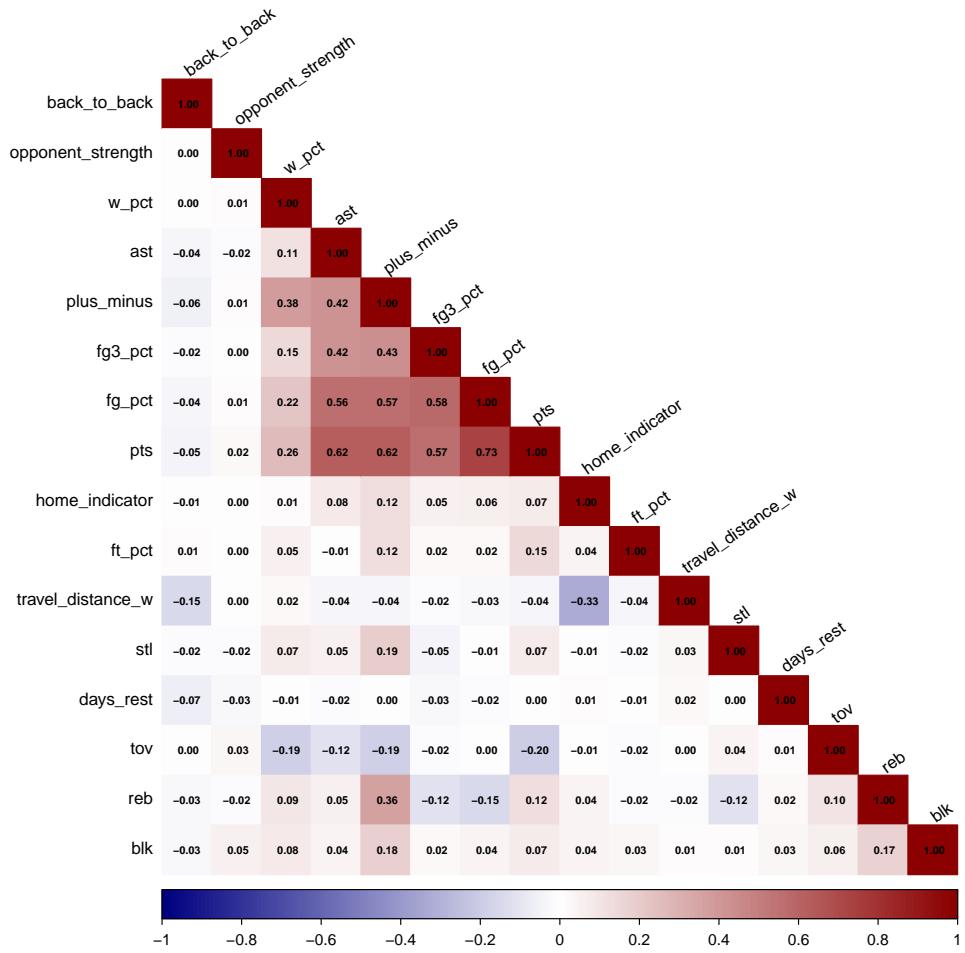
## corrplot 0.95 loaded

# ---- Select only important model features ----
important_vars <- nba %>%
  select(
    w_pct,                      # target
    travel_distance_w,          # travel
    days_rest,
    back_to_back,
    home_indicator,
    opponent_strength,         # difficulty factor
    fg_pct, fg3_pct, ft_pct,   # shooting efficiency
    reb, ast, stl, blk, tov,   # core box score stats
    plus_minus, pts            # scoring metrics
  )

# --- Ensure all variables are numeric ---
important_vars <- important_vars %>%
  mutate(across(everything(), as.numeric))

# --- Compute correlation matrix ---
cor_mat <- cor(important_vars, use = "pairwise.complete.obs")

corrplot(
  round(cor_mat, 2),
  method = "color",
  type = "lower",
  order = "hclust",
  tl.col = "black",
  tl.srt = 35,
  tl.cex = 0.9,
  number.cex = 0.55,
  number.font = 2,
  addCoef.col = "black",
  mar = c(1,1,1,1),
  cl.cex = 0.8,
  col = colorRampPalette(c("navy", "white", "darkred"))(200)
)
```



## Modeling

```
# Prepare data for modeling
nba_glm <- nba %>%
  mutate(
    win = ifelse(wl == "W", 1, 0),
    back_to_back_num = ifelse(back_to_back == "1", 1, 0),
    travel_distance_w = ifelse(is.na(travel_distance_w), travel_distance, travel_distance_w),
    home_indicator = ifelse(home_away == "Home", 1, 0)
  ) %>%
  select(
    win,
    # Schedule factors
    travel_distance_w,
    days_rest,
```

```

back_to_back_num,
home_indicator,
opponent_strength,
# ADD safe performance variables (do NOT leak outcome)
fg_pct,
fg3_pct,
ast,
reb,
tov
) %>%
tidyverse::drop_na()

```

## Logistic Regression Model

```

full_glm <- glm(
  win ~ travel_distance_w + days_rest + back_to_back_num +
    home_indicator + opponent_strength +
    fg_pct + fg3_pct + ast + reb + tov,
  data = nba_glm,
  family = binomial(link = "logit")
)

summary(full_glm)

## 
## Call:
## glm(formula = win ~ travel_distance_w + days_rest + back_to_back_num +
##       home_indicator + opponent_strength + fg_pct + fg3_pct + ast +
##       reb + tov, family = binomial(link = "logit"), data = nba_glm)
## 
## Coefficients:
##                               Estimate Std. Error z value Pr(>|z|)
## (Intercept)           -2.396e+01  7.351e-01 -32.597 < 2e-16 ***
## travel_distance_w   5.539e-05  7.705e-05   0.719 0.472219
## days_rest            1.215e-03  2.486e-03   0.489 0.624959
## back_to_back_num   -1.133e-01  1.023e-01  -1.108 0.267958
## home_indicator      3.038e-01  8.109e-02   3.746 0.000180 ***
## opponent_strength   4.938e-01  3.391e-01   1.456 0.145381
## fg_pct              3.127e+01  1.201e+00  26.035 < 2e-16 ***
## fg3_pct              6.643e+00  6.069e-01  10.946 < 2e-16 ***
## ast                 -3.524e-02  9.365e-03  -3.763 0.000168 ***
## reb                 2.233e-01  7.970e-03  28.023 < 2e-16 ***
## tov                 -1.707e-01  1.074e-02 -15.901 < 2e-16 ***
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## (Dispersion parameter for binomial family taken to be 1)
## 
## Null deviance: 6820.6 on 4919 degrees of freedom
## Residual deviance: 4188.3 on 4909 degrees of freedom

```

```

## AIC: 4210.3
##
## Number of Fisher Scoring iterations: 5

ll_null <- logLik(glm(win ~ 1, data = nba_glm, family = binomial))
ll_full <- logLik(full_glm)

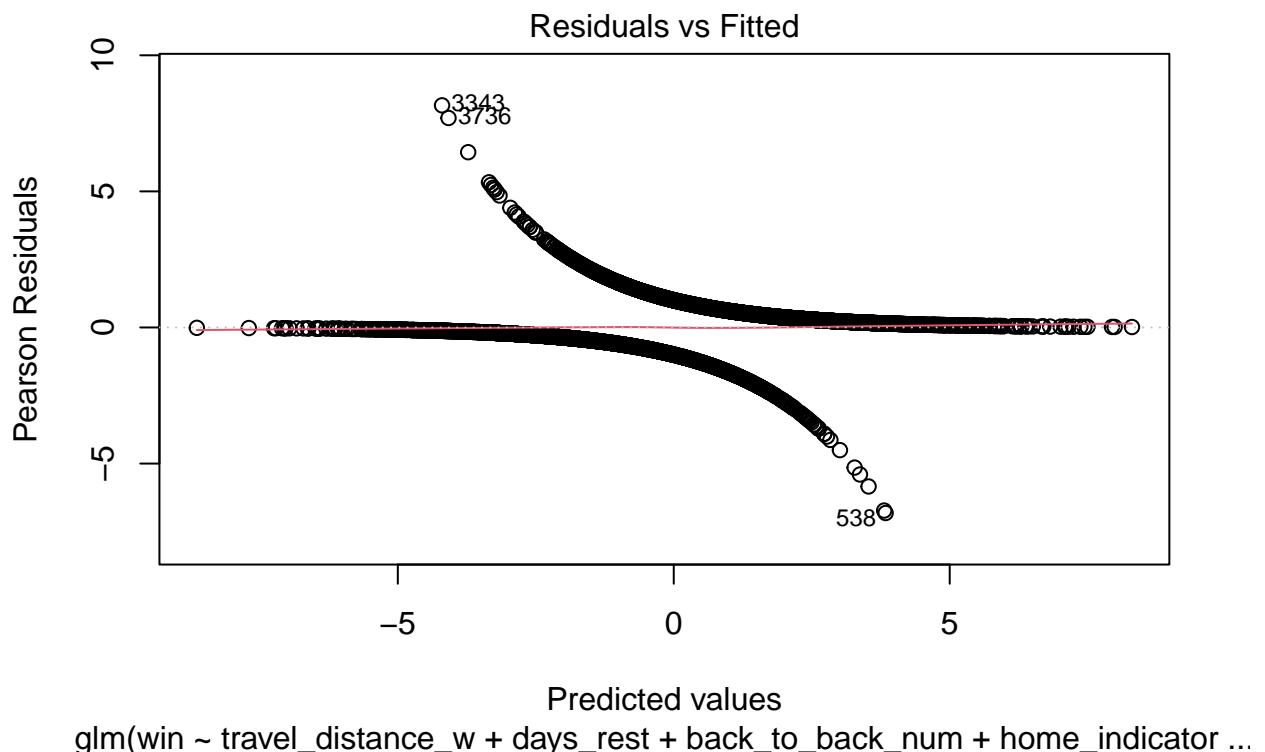
# McFadden R2
R2 <- 1 - (as.numeric(ll_full) / as.numeric(ll_null))
R2

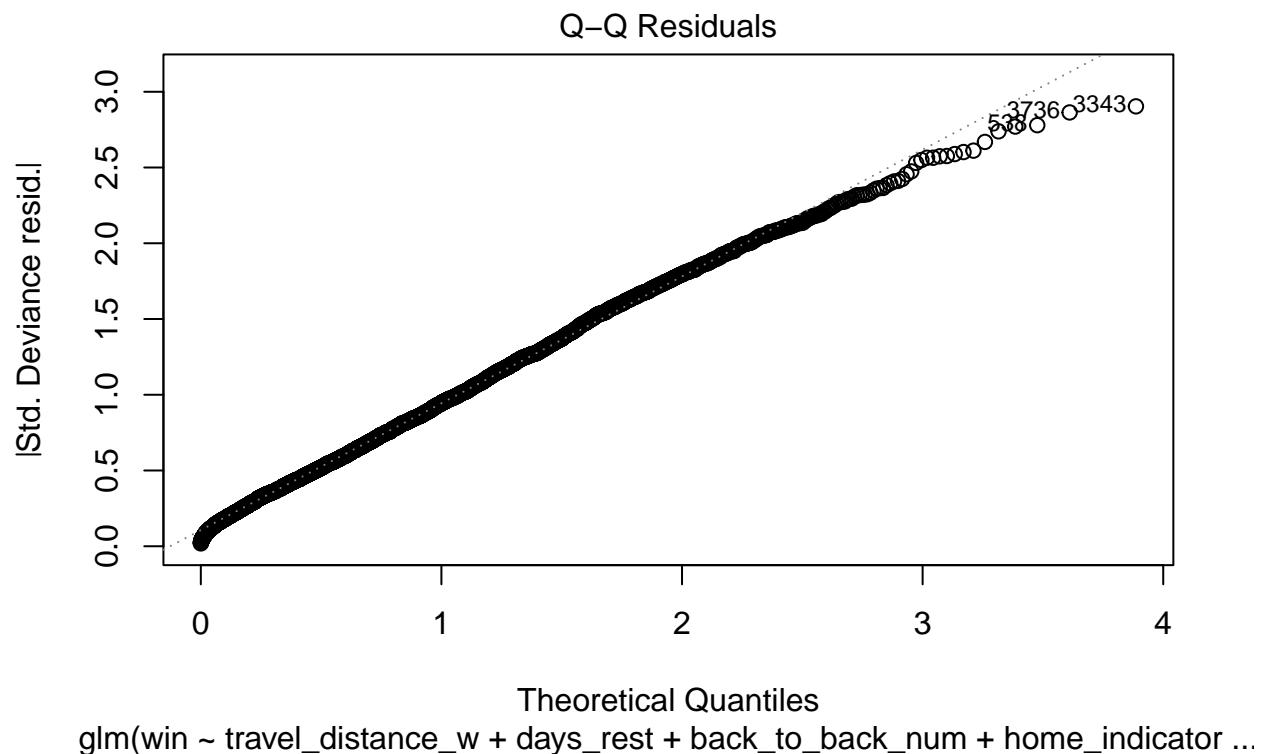
## [1] 0.38593

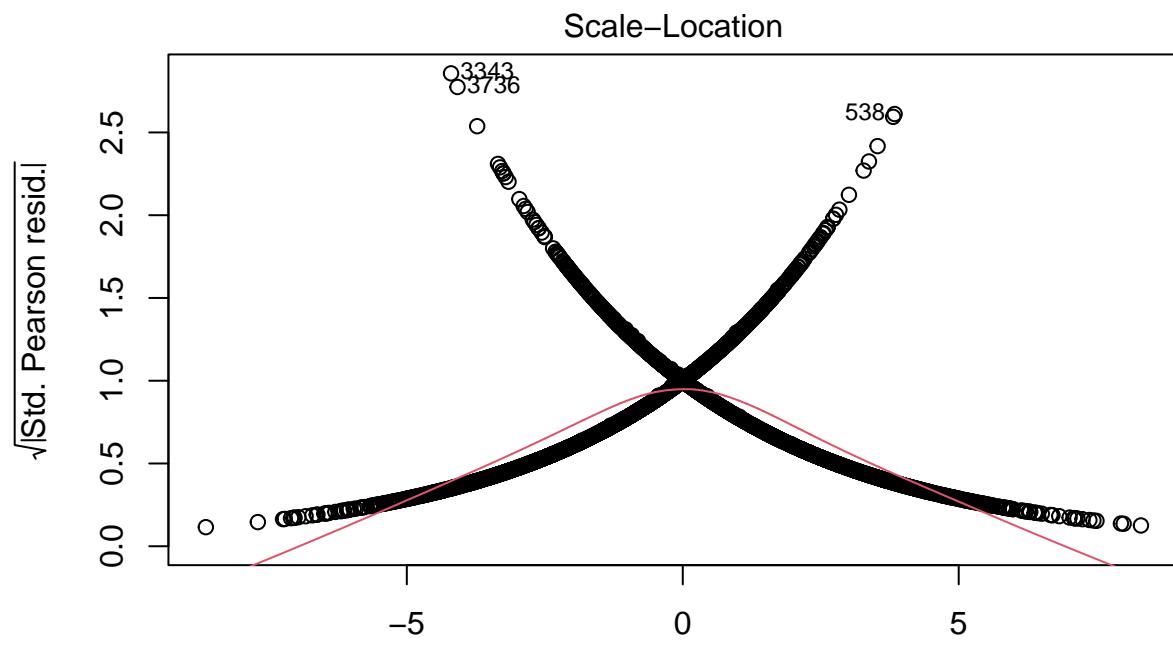
```

## Marginal Plots

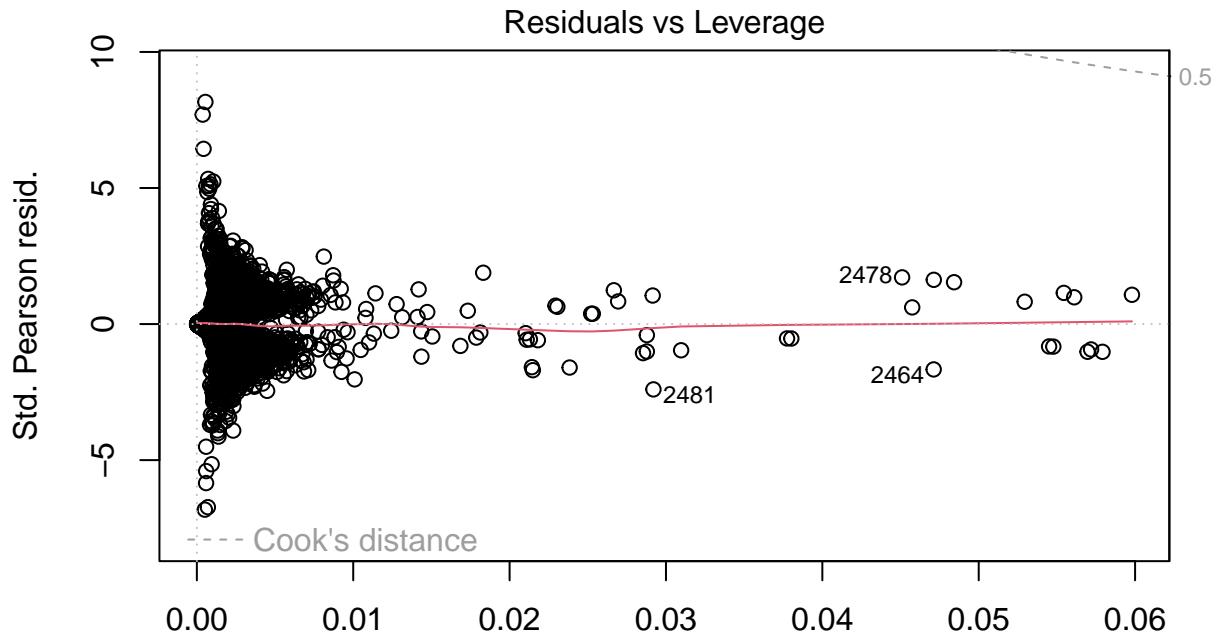
```
plot(full_glm)
```







Predicted values  
glm(win ~ travel\_distance\_w + days\_rest + back\_to\_back\_num + home\_indicator ...)



Leverage  
`glm(win ~ travel_distance_w + days_rest + back_to_back_num + home_indicator ...)`

```
library(ggplot2)
library(patchwork)

df <- model.frame(full_glm)

make_typical <- function(df) {
  out <- df[1, , drop = FALSE]
  for (nm in names(df)) {
    if (is.numeric(df[[nm]])) {
      out[[nm]] <- base::mean(df[[nm]], na.rm = TRUE)
    } else {
      out[[nm]] <- levels(df[[nm]])[1]
    }
  }
  out
}

marginal_plot <- function(model, df, var, label = var, n = 100) {
  base_row <- make_typical(df)

  if (is.numeric(df[[var]])) {
    lo <- stats::quantile(df[[var]], 0.05, na.rm = TRUE)
    hi <- stats::quantile(df[[var]], 0.95, na.rm = TRUE)
    grid <- base_row[rep(1, n), , drop = FALSE]
    grid[[var]] <- seq(lo, hi, length.out = n)
  } else {
```

```

levs <- levels(factor(df[[var]]))
grid <- do.call(rbind, lapply(levs, function(L) {
  r <- base_row
  r[[var]] <- factor(L, levels = levs)
  r
})))
}

p <- suppressWarnings(predict(model, newdata = grid, type = "link", se.fit = TRUE))
grid$fit <- plogis(p$fit)
grid$lwr <- plogis(p$fit - 1.96 * p$se.fit)
grid$upr <- plogis(p$fit + 1.96 * p$se.fit)

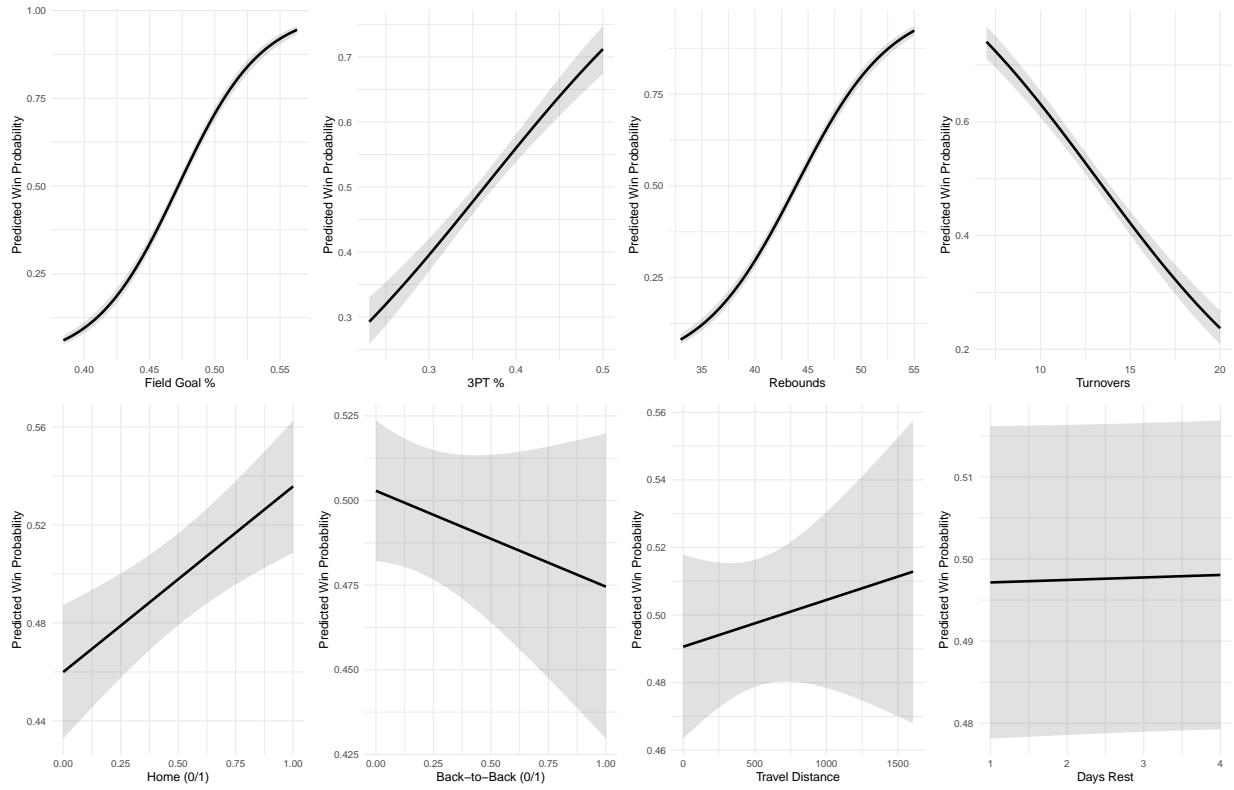
g <- if (is.numeric(df[[var]])) {
  ggplot(grid, aes(x = .data[[var]], y = fit)) +
    geom_ribbon(aes(ymin = lwr, ymax = upr), alpha = 0.15) +
    geom_line(linewidth = 1)
} else {
  ggplot(grid, aes(x = .data[[var]], y = fit)) +
    geom_point(size = 3) +
    geom_errorbar(aes(ymin = lwr, ymax = upr), width = 0.15)
}

g +
  labs(x = label, y = "Predicted Win Probability") +
  theme_minimal(base_size = 10) +
  theme(
    plot.title = element_blank(), # ensure no title
    plot.margin = margin(6, 6, 6, 6)
  )
}

p1 <- marginal_plot(full_glm, df, "fg_pct", "Field Goal %")
p2 <- marginal_plot(full_glm, df, "fg3_pct", "3PT %")
p3 <- marginal_plot(full_glm, df, "reb", "Rebounds")
p4 <- marginal_plot(full_glm, df, "tov", "Turnovers")
p5 <- marginal_plot(full_glm, df, "home_indicator", "Home (0/1)")
p6 <- marginal_plot(full_glm, df, "back_to_back_num", "Back-to-Back (0/1)")
p7 <- marginal_plot(full_glm, df, "travel_distance_w", "Travel Distance")
p8 <- marginal_plot(full_glm, df, "days_rest", "Days Rest")

# Combine 2x4 with no global annotation
(p1 | p2 | p3 | p4) /
(p5 | p6 | p7 | p8)

```



```
## BIC, AIC
```

```
# Backward elimination (step using AIC; start from full model)
# (uses base step() which by default uses AIC; change k=log(n) for BIC)
backward_aic <- step(full_glm, direction = "backward", trace = 1)
```

```
## Start: AIC=4210.31
## win ~ travel_distance_w + days_rest + back_to_back_num + home_indicator +
##       opponent_strength + fg_pct + fg3_pct + ast + reb + tov
##
##                               Df Deviance    AIC
## - days_rest             1   4188.5 4208.5
## - travel_distance_w    1   4188.8 4208.8
## - back_to_back_num      1   4189.5 4209.5
## <none>                  4188.3 4210.3
## - opponent_strength     1   4190.4 4210.4
## - home_indicator        1   4202.4 4222.4
## - ast                    1   4202.6 4222.6
## - fg3_pct                1   4314.7 4334.7
## - tov                    1   4470.4 4490.4
## - fg_pct                 1   5122.0 5142.0
## - reb                    1   5365.9 5385.9
##
## Step: AIC=4208.54
## win ~ travel_distance_w + back_to_back_num + home_indicator +
##       opponent_strength + fg_pct + fg3_pct + ast + reb + tov
##
##                               Df Deviance    AIC
```

```

## - travel_distance_w 1 4189.1 4207.1
## - back_to_back_num 1 4189.9 4207.9
## <none> 4188.5 4208.5
## - opponent_strength 1 4190.6 4208.6
## - home_indicator 1 4202.6 4220.6
## - ast 1 4202.9 4220.9
## - fg3_pct 1 4314.7 4332.7
## - tov 1 4470.5 4488.5
## - fg_pct 1 5122.2 5140.2
## - reb 1 5366.6 5384.6
##
## Step: AIC=4207.07
## win ~ back_to_back_num + home_indicator + opponent_strength +
##      fg_pct + fg3_pct + ast + reb + tov
##
##          Df Deviance    AIC
## - back_to_back_num 1 4190.7 4206.7
## <none> 4189.1 4207.1
## - opponent_strength 1 4191.2 4207.2
## - home_indicator 1 4202.9 4218.9
## - ast 1 4203.4 4219.4
## - fg3_pct 1 4315.1 4331.1
## - tov 1 4471.2 4487.2
## - fg_pct 1 5122.3 5138.3
## - reb 1 5366.6 5382.6
##
## Step: AIC=4206.71
## win ~ home_indicator + opponent_strength + fg_pct + fg3_pct +
##      ast + reb + tov
##
##          Df Deviance    AIC
## <none> 4190.7 4206.7
## - opponent_strength 1 4192.8 4206.8
## - home_indicator 1 4204.5 4218.5
## - ast 1 4205.0 4219.0
## - fg3_pct 1 4316.5 4330.5
## - tov 1 4473.1 4487.1
## - fg_pct 1 5129.1 5143.1
## - reb 1 5372.9 5386.9

summary(backward_aic)

```

```

##
## Call:
## glm(formula = win ~ home_indicator + opponent_strength + fg_pct +
##      fg3_pct + ast + reb + tov, family = binomial(link = "logit"),
##      data = nba_glm)
##
## Coefficients:
##                               Estimate Std. Error z value Pr(>|z|)
## (Intercept)           -23.957177   0.729446 -32.843 < 2e-16 ***
## home_indicator        0.285529   0.076868   3.715 0.000204 ***
## opponent_strength     0.491201   0.338708   1.450 0.146997
## fg_pct                31.313648   1.200284  26.089 < 2e-16 ***

```

```

## fg3_pct      6.622838  0.606258 10.924 < 2e-16 ***
## ast         -0.035225  0.009355 -3.765 0.000166 ***
## reb          0.223482  0.007962 28.068 < 2e-16 ***
## tov         -0.170740  0.010732 -15.910 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 6820.6 on 4919 degrees of freedom
## Residual deviance: 4190.7 on 4912 degrees of freedom
## AIC: 4206.7
##
## Number of Fisher Scoring iterations: 5

# --- Forward selection (start from intercept; use AIC) ---
null_mod <- glm(win ~ 1, data = nba_glm) # start with intercept only
scope_formula <- formula(full_glm)

forward_aic <- step(null_mod, scope = scope_formula, direction = "forward", trace = 1)

## Start: AIC=7145.79
## win ~ 1
##
##                               Df Deviance   AIC
## + fg_pct                  1   958.96 5923.1
## + fg3_pct                 1  1078.37 6500.5
## + ast                     1  1106.15 6625.6
## + reb                     1  1122.69 6698.6
## + tov                     1  1204.28 7043.8
## + home_indicator          1  1220.69 7110.4
## + back_to_back_num        1  1226.65 7134.4
## + travel_distance_w       1  1229.00 7143.8
## <none>                   1230.00 7145.8
## + opponent_strength       1  1229.95 7147.6
## + days_rest               1  1229.99 7147.7
##
## Step: AIC=5923.11
## win ~ fg_pct
##
##                               Df Deviance   AIC
## + reb                      1   789.03 4965.5
## + tov                      1   933.78 5794.2
## + fg3_pct                  1   947.04 5863.6
## + ast                      1   953.75 5898.3
## + home_indicator           1   954.59 5902.6
## + back_to_back_num         1   957.76 5919.0
## <none>                    958.96 5923.1
## + travel_distance_w        1   958.73 5923.9
## + days_rest                1   958.90 5924.8
## + opponent_strength        1   958.96 5925.1
##
## Step: AIC=4965.47
## win ~ fg_pct + reb

```

```

##                                     Df Deviance    AIC
## + tov                         1   748.19 4706.0
## + fg3_pct                      1   773.30 4868.4
## + home_indicator                1   787.11 4955.5
## + back_to_back_num              1   788.64 4965.1
## <none>                         789.03 4965.5
## + opponent_strength             1   788.92 4966.8
## + ast                           1   788.99 4967.2
## + travel_distance_w             1   789.01 4967.3
## + days_rest                     1   789.03 4967.5
##
## Step:  AIC=4705.99
## win ~ fg_pct + reb + tov
##
##                                     Df Deviance    AIC
## + fg3_pct                      1   733.20 4608.4
## + home_indicator                1   746.48 4696.8
## + ast                           1   747.49 4703.4
## + back_to_back_num              1   747.83 4705.6
## <none>                         748.19 4706.0
## + opponent_strength             1   747.93 4706.3
## + travel_distance_w             1   748.17 4707.9
## + days_rest                     1   748.18 4707.9
##
## Step:  AIC=4608.41
## win ~ fg_pct + reb + tov + fg3_pct
##
##                                     Df Deviance    AIC
## + ast                           1   731.04 4595.9
## + home_indicator                1   731.70 4600.3
## + back_to_back_num              1   732.81 4607.8
## <none>                         733.20 4608.4
## + opponent_strength             1   732.91 4608.5
## + days_rest                     1   733.17 4610.2
## + travel_distance_w             1   733.19 4610.4
##
## Step:  AIC=4595.92
## win ~ fg_pct + reb + tov + fg3_pct + ast
##
##                                     Df Deviance    AIC
## + home_indicator                1   729.38 4586.7
## + back_to_back_num              1   730.63 4595.2
## <none>                         731.04 4595.9
## + opponent_strength             1   730.80 4596.3
## + days_rest                     1   731.02 4597.8
## + travel_distance_w             1   731.03 4597.8
##
## Step:  AIC=4586.73
## win ~ fg_pct + reb + tov + fg3_pct + ast + home_indicator
##
##                                     Df Deviance    AIC
## + back_to_back_num              1   728.99 4586.1
## <none>                         729.38 4586.7

```

```

## + opponent_strength 1 729.13 4587.1
## + travel_distance_w 1 729.28 4588.0
## + days_rest 1 729.36 4588.6
##
## Step: AIC=4586.07
## win ~ fg_pct + reb + tov + fg3_pct + ast + home_indicator + back_to_back_num
##
##          Df Deviance    AIC
## <none>      728.99 4586.1
## + opponent_strength 1 728.74 4586.4
## + travel_distance_w 1 728.94 4587.7
## + days_rest 1 728.98 4588.0

```

```
summary(forward_aic)
```

```

##
## Call:
## glm(formula = win ~ fg_pct + reb + tov + fg3_pct + ast + home_indicator +
##       back_to_back_num, data = nba_glm)
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2.7543313  0.0667508 -41.263 < 2e-16 ***
## fg_pct        4.3485492  0.1383656  31.428 < 2e-16 ***
## reb           0.0304262  0.0008563  35.531 < 2e-16 ***
## tov           -0.0242409  0.0014415 -16.816 < 2e-16 ***
## fg3_pct        0.8711476  0.0830518  10.489 < 2e-16 ***
## ast           -0.0052743  0.0013290  -3.969 7.34e-05 ***
## home_indicator 0.0367596  0.0110334   3.332  0.00087 ***
## back_to_back_num -0.0234947  0.0144079  -1.631  0.10302
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.1484095)
##
## Null deviance: 1230.00 on 4919 degrees of freedom
## Residual deviance: 728.99 on 4912 degrees of freedom
## AIC: 4586.1
##
## Number of Fisher Scoring iterations: 2

```

## Check for Multicollinearity

```
# VIF Check
library(car)
```

```

## Loading required package: carData

##
## Attaching package: 'car'
```

```

## The following object is masked from 'package:DescTools':
##
##      Recode

## The following object is masked from 'package:dplyr':
##
##      recode

## The following object is masked from 'package:purrr':
##
##      some

vif_values <- vif(full_glm)
vif_values

```

	days_rest	back_to_back_num	home_indicator
## travel_distance_w	1.140456	1.008810	1.116124
## opponent_strength		fg_pct	fg3_pct
##	1.004421	1.877146	ast
##	reb	tov	
##	1.608652	1.147233	1.320447

The model performs poorly on held-out testing data even VIF values show  $< 2$ . Classic symptom of original overfitting. Then we need Lasso/Ridge.

## Train/Test Split

```

set.seed(123)

n <- nrow(nba_glm)
train_idx <- sample(1:n, size = 0.7 * n)

train_data <- nba_glm[train_idx, ]
test_data <- nba_glm[-train_idx, ]

y_train <- train_data$win
y_test <- test_data$win

# Create the model
X_train <- model.matrix(win ~ ., data = train_data)[, -1]
X_test <- model.matrix(win ~ ., data = test_data)[, -1]

```

```
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
## Attaching package: 'pROC'
```

```

## The following objects are masked from 'package:stats':
##
##      cov, smooth, var

# Convert X_test to data frame (glm requires data.frame)
X_test <- as.data.frame(X_test)

# 1. Get predicted probabilities
log_prob <- predict(full_glm, newdata = X_test, type = "response")

# 2. Convert y_test to numeric (for WIN target)
y_num <- as.numeric(as.character(y_test))

# 3. AUC
log_auc <- auc(y_num, log_prob)

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

log_auc

## Area under the curve: 0.8766

# 4. Convert to classes
log_pred <- ifelse(log_prob > 0.5, 1, 0)

# 5. Accuracy
log_accuracy <- base::mean(log_pred == y_num)
log_accuracy

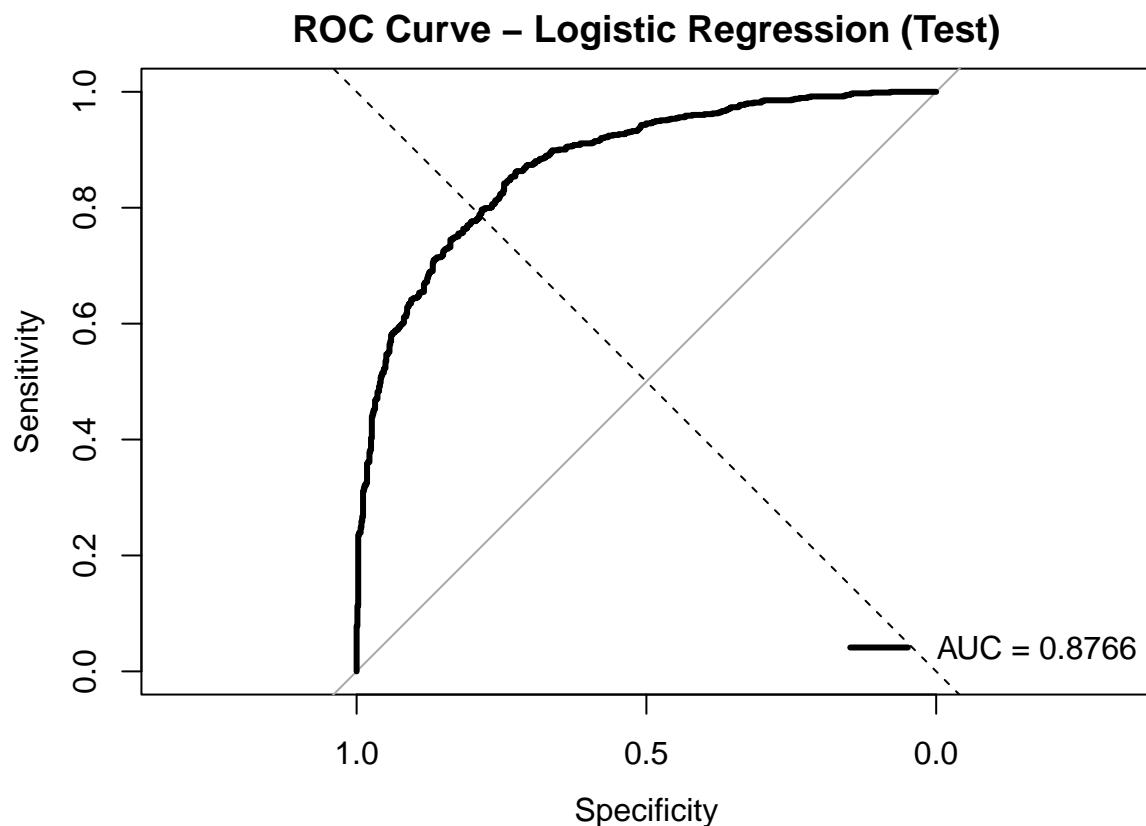
## [1] 0.7879404

library(pROC)

log_roc <- roc(response = y_num, predictor = log_prob, levels = c(0,1), direction = "<")

# plot
plot(log_roc, lwd = 3, main = "ROC Curve - Logistic Regression (Test)")
abline(a = 0, b = 1, lty = 2) # diagonal
legend("bottomright",
       legend = paste0("AUC = ", round(auc(log_roc), 4)),
       lwd = 3, bty = "n")

```



```

library(caret)

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following objects are masked from 'package:DescTools':
##      MAE, RMSE

## The following object is masked from 'package:purrr':
##      lift

library(ggplot2)
library(dplyr)

# --- Confusion matrix (with metrics) ---
ref <- factor(y_num, levels = c(0, 1)) # actual
pred <- factor(log_pred, levels = c(0, 1)) # predicted

cm <- caret::confusionMatrix(pred, ref, positive = "1")
cm

```

```

## Confusion Matrix and Statistics
##
##             Reference
## Prediction  0   1
##           0 567 157
##           1 156 596
##
##                 Accuracy : 0.7879
##                   95% CI : (0.7662, 0.8085)
##       No Information Rate : 0.5102
##     P-Value [Acc > NIR] : <2e-16
##
##                 Kappa : 0.5757
##
## McNemar's Test P-Value : 1
##
##                 Sensitivity : 0.7915
##                 Specificity : 0.7842
##      Pos Pred Value : 0.7926
##      Neg Pred Value : 0.7831
##          Prevalence : 0.5102
##      Detection Rate : 0.4038
## Detection Prevalence : 0.5095
##     Balanced Accuracy : 0.7879
##
##     'Positive' Class : 1
##

# Quick access to key metrics:
cm$overall["Accuracy"]

## Accuracy
## 0.7879404

cm$byClass[c("Sensitivity", "Specificity", "Balanced Accuracy")]

##      Sensitivity      Specificity Balanced Accuracy
## 0.7915007        0.7842324        0.7878665

# --- Plot heatmap of confusion matrix ---
cm_df <- as.data.frame(cm$table) %>%
  mutate(
    Reference = factor(Reference, levels = c("0", "1"), labels = c("Actual: Loss", "Actual: Win")),
    Prediction = factor(Prediction, levels = c("0", "1"), labels = c("Pred: Loss", "Pred: Win"))
  )

ggplot(cm_df, aes(x = Reference, y = Prediction, fill = Freq)) +
  geom_tile() +
  geom_text(aes(label = Freq), size = 5, fontface = "bold") +
  scale_fill_gradient(low = "white", high = "steelblue") +
  labs(
    title = "Confusion Matrix - Logistic Regression (Test)",

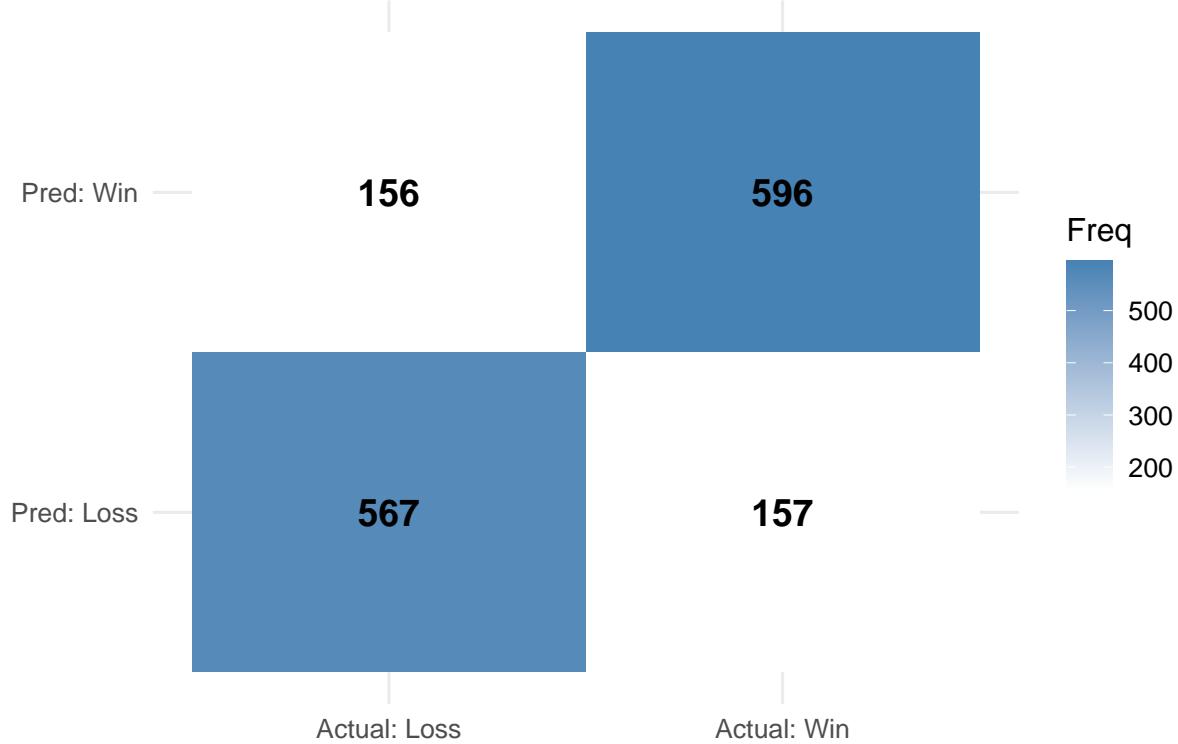
```

```

  x = "", y =
) +
theme_minimal(base_size = 12)

```

Confusion Matrix – Logistic Regression (Test)



## Ridge and Lasso Logistic Regression

```

# Ridge (alpha=0)
library(glmnet)

set.seed(123)

ridge_cv <- cv.glmnet(
  X_train, y_train,
  family = "binomial",
  alpha = 0
)

ridge_model <- glmnet(
  X_train, y_train,
  family = "binomial",
  alpha = 0,
  lambda = ridge_cv$lambda.min
)

```

```

)

# Lasso (alpha=1)
set.seed(123)

lasso_cv <- cv.glmnet(
  X_train, y_train,
  family = "binomial",
  alpha = 1
)

lasso_model <- glmnet(
  X_train, y_train,
  family = "binomial",
  alpha = 1,
  lambda = lasso_cv$lambda.min
)

# 1) Capture the training column order
train_cols <- colnames(X_train)

# 2) Recreate the test design matrix with the SAME formula you used for X_train
#     If you built X_train via: model.matrix(win ~ ., data = train_data)[,-1]
#     then do the same here:
X_test <- model.matrix(win ~ ., data = test_data)[, -1, drop = FALSE]

# 3) Add any columns that exist in train but not in test (rare, e.g., missing factor levels)
missing_cols <- setdiff(train_cols, colnames(X_test))
if (length(missing_cols) > 0) {
  add_mat <- matrix(0, nrow = nrow(X_test), ncol = length(missing_cols))
  colnames(add_mat) <- missing_cols
  X_test <- cbind(X_test, add_mat)
}

# 4) Drop any extra columns that appear in test but not train, and reorder to match train
X_test <- X_test[, train_cols, drop = FALSE]

# 5) Ensure plain numeric matrix (glmnet calls as.matrix internally)
storage.mode(X_test) <- "double"

```

## Predict on TEST DATA

```

# Ridge
ridge_prob <- as.vector(predict(ridge_model, newx = X_test, type = "response"))

# Lasso
lasso_prob <- as.vector(predict(lasso_model, newx = X_test, type = "response"))

```

## Confusion Matrix

```
## ---- confusion_matrix_plots, message=FALSE, warning=FALSE -----

# Ensure predicted classes exist
ridge_pred <- ifelse(ridge_prob >= 0.5, 1, 0)
lasso_pred <- ifelse(lasso_prob >= 0.5, 1, 0)

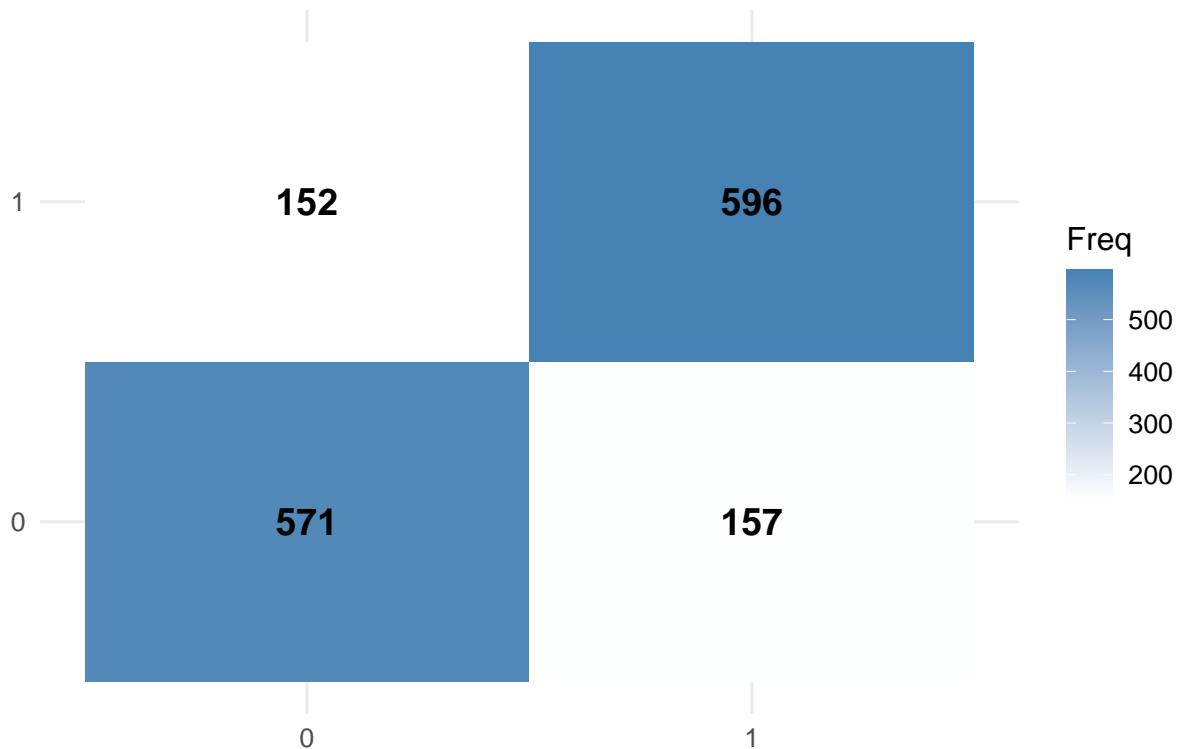
# Helper to build cm data.frame in the shape you want
make_cm_df <- function(actual, predicted) {
  cm <- table(Reference = actual, Prediction = predicted)
  df <- as.data.frame(cm)
  # enforce 0/1 order for clean axis display
  df$Reference <- factor(df$Reference, levels = c(0,1))
  df$Prediction <- factor(df$Prediction, levels = c(0,1))
  df
}

cm_df_ridge <- make_cm_df(y_test, ridge_pred)
cm_df_lasso <- make_cm_df(y_test, lasso_pred)

library(ggplot2)

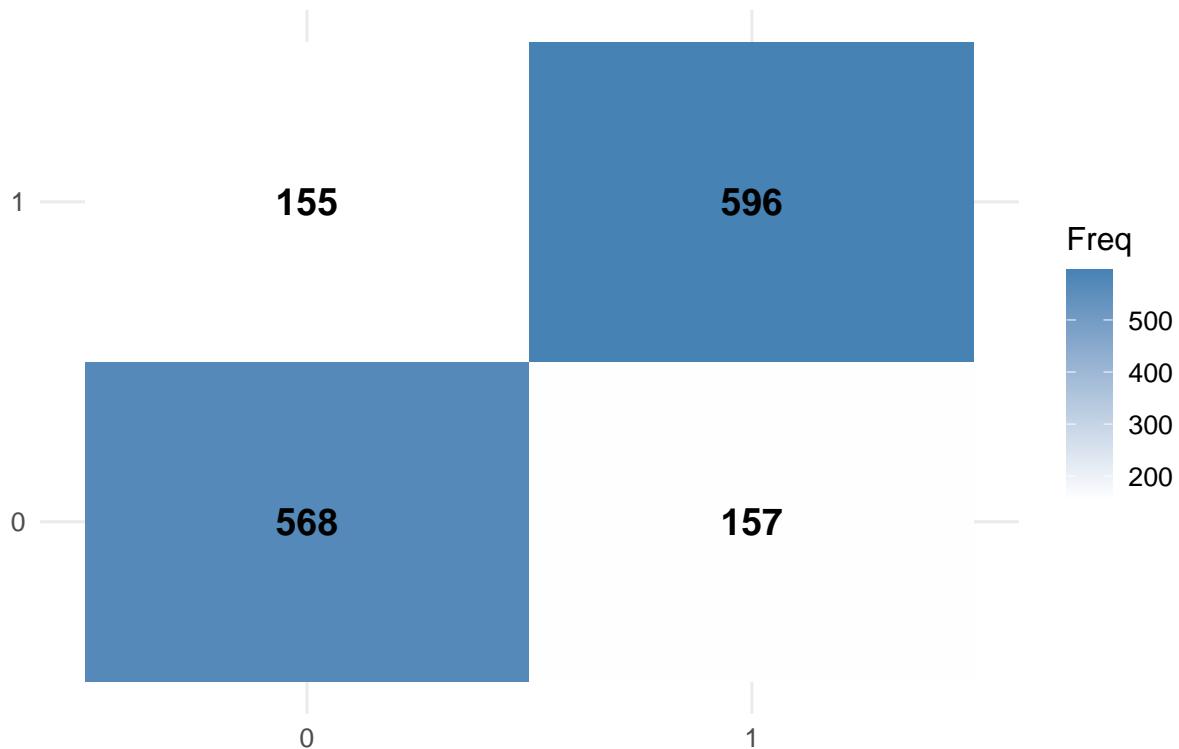
# Plot: Ridge
ggplot(cm_df_ridge, aes(x = Reference, y = Prediction, fill = Freq)) +
  geom_tile() +
  geom_text(aes(label = Freq), size = 5, fontface = "bold") +
  scale_fill_gradient(low = "white", high = "steelblue") +
  labs(
    title = "Confusion Matrix - Ridge (Test)",
    x = "", y = ""
  ) +
  theme_minimal(base_size = 12)
```

Confusion Matrix – Ridge (Test)



```
# Plot: Lasso
ggplot(cm_df_lasso, aes(x = Reference, y = Prediction, fill = Freq)) +
  geom_tile() +
  geom_text(aes(label = Freq), size = 5, fontface = "bold") +
  scale_fill_gradient(low = "white", high = "steelblue") +
  labs(
    title = "Confusion Matrix - Lasso (Test)",
    x = "", y = ""
  ) +
  theme_minimal(base_size = 12)
```

Confusion Matrix – Lasso (Test)



### Compute AUC, Accuracy and ROC Curves

```
library(pROC)

# AUC
ridge_auc <- auc(y_test, ridge_prob)

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

lasso_auc <- auc(y_test, lasso_prob)

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

ridge_auc

## Area under the curve: 0.8734
```

```

lasso_auc

## Area under the curve: 0.8758

# Convert to classes
ridge_pred <- ifelse(ridge_prob > 0.5, 1, 0)
lasso_pred <- ifelse(lasso_prob > 0.5, 1, 0)

# Accuracy
ridge_accuracy <- base::mean(ridge_pred == y_test)
lasso_accuracy <- base::mean(lasso_pred == y_test)

ridge_accuracy

```

## [1] 0.7906504

```

lasso_accuracy

## [1] 0.7886179

```

## Plot ROC Curves

```

ridge_roc <- roc(response = y_test,
                    predictor = ridge_prob,
                    levels = c(0, 1),
                    direction = "<",
                    quiet = TRUE)

lasso_roc <- roc(response = y_test,
                    predictor = lasso_prob,
                    levels = c(0, 1),
                    direction = "<",
                    quiet = TRUE)

plot(ridge_roc, col = "darkred", lwd = 3,
      main = "ROC Curve: Ridge vs Lasso (Test Data)")

plot(lasso_roc, col = "darkgreen", lwd = 3, add = TRUE)

legend("bottomright",
       legend = c(
         paste0("Ridge AUC = ", round(ridge_auc, 3)),
         paste0("Lasso AUC = ", round(lasso_auc, 3))
       ),
       col = c("darkred", "darkgreen"),
       lwd = 3)

```

**ROC Curve: Ridge vs Lasso (Test Data)**

